

BenGER: Benchmarking LLM Systems on Subsumption-Based Legal Reasoning in German Law

Sebastian Nagl^{1*} Ann-Kristin Mayrhofer^{2†} Martin Heidebach^{2†}
Aleyna Koçak^{3†} Anne Zettelmeier^{4†} Elly Breu^{1†}
Angelina Greiner^{1†} Sofija Milijas^{1†} Matthias Grabmair¹

¹Technical University of Munich (TUM) ²Ludwig Maximilian University of Munich (LMU)

³University of Konstanz ⁴University of Saarbrücken

Abstract

We introduce BenGER (Benchmark for German Law), a benchmark and dataset for subsumption-based legal reasoning in German law. It combines 596 exam-style free-text case tasks across several levels of legal education with 531 short doctrinal reasoning tasks, and a controlled subset adds timed human solutions written both unaided and in human-AI co-creation. We evaluate 12 contemporary LLM systems (closed flagship, efficiency-oriented, and open-weight) with a rubric-aligned LLM-as-a-Judge, which we cross-validate against a multi-rater human layer (three blind reviews per solution) and across six LLM judges. Our results show that closed-flagship systems lead across all three corpora, that human-AI co-creation measurably improves on unaided human work, and that the judge tracks human grading at Pearson $r=0.76$ and Cohen’s $\kappa=0.60$, with rankings stable across judges and two judges from independent providers clearing the Calderon single-reviewer replacement bar on human-written solutions.

1 Introduction

Evaluating legal reasoning is hard. Even human experts grade the same answer very differently. Hufeld (Hufeld, 2024) shows that identical answers in German law exams receive markedly different scores depending on the grader. Doctrinal analysis is interpretive, so this is no surprise. But it is a problem for benchmarking large language models (LLMs). If human evaluation is itself noisy, what does it mean to benchmark legal reasoning?

Legal NLP benchmarks have evolved fast. They moved from classification (Chalkidis et al., 2022; Guha et al., 2023) to reasoning (Fei et al., 2024; Fan et al., 2026; Jia et al., 2026; Shi et al., 2026) and retrieval-augmented evaluation (Pipitone and

Houir Alami, 2024; Butler and Butler, 2026). But most still emphasise short-form QA, common-law jurisdictions, or stable ground-truth labels. Open-ended subsumption-based German legal reasoning under noisy expert judgement remains underserved.

We address these gaps with a benchmark centered on subsumption-based legal reasoning, the core reasoning paradigm in German law (Section 2) and, more broadly, in many civil-law jurisdictions worldwide.

We present the BenGER (Benchmark for German Law) dataset¹ for evaluating LLMs on subsumption-based legal reasoning in German law. It has three components. Together they cover 596 exam-style free-text case tasks across different stages of legal education and 531 short doctrinal reasoning tasks. It also adds human baselines, both traditional legal work and human-AI co-creation, on a controlled Benchathon validation subset. Figure 1 summarises the three components and the generation-plus-evaluation pipeline. Generation, annotation, and evaluation ran end-to-end on the BenGER platform (Nagl and Grabmair, 2026), the open-source web application we built earlier for benchmarking legal-domain LLM systems.

How do we evaluate free-text reasoning when the human judgement is itself noisy? We adopt a hybrid framework. An LLM-as-a-Judge, aligned with German grading rubrics, does the scoring. We cross-validate it against three blind human reviews per solution, plus an un-blind review by the task’s creator as a reference signal. We treat evaluation as a distribution over plausible expert assessments, not a single deterministic label.

We make three contributions. (i) A benchmark for subsumption-based legal reasoning in German law. (ii) An empirical evaluation of 12 contem-

* Core Contributor, corresponding author. Correspondence to sebastian.nagl@tum.de.

† Dataset contributor.

¹Full code and dataset are available at <https://github.com/SebastianNagl/benger-platform>.




Figure 1: Overview of the benchmark and evaluation pipeline. Three German-law corpora (the published ZJS exam corpus, the short Doctrinal Principles question set, and the controlled Benchathon validation subset with human solutions in two working conditions) are answered by a panel of LLM systems and scored by a rubric-aligned LLM judge. The Benchathon subset is additionally graded by a blind human reviewer pool (with one creator review per solution); the resulting human distribution is used to validate the LLM judge.

porary LLM systems (closed flagship, efficiency-oriented, and open-weight) alongside human and co-creative baselines. (iii) An evaluation methodology that accounts for the inherent variability of legal grading.

We organise the empirical analysis of the full BenGER dataset around five questions. **RQ1 (Performance):** how well do contemporary LLM systems solve open-ended subsumption-based German legal reasoning tasks, and how do they compare to human baselines? **RQ2 (System-tier effects):** how does performance vary across system tiers and open-/closed-weight access modes? **RQ3 (Evaluation validity):** how well do generic automatic metrics correlate with rubric-based legal evaluation? **RQ4 (Human-AI co-creation):** does AI-assisted drafting improve human performance, and how does it compare to standalone LLM systems? **RQ5 (Grading reliability):** how closely do LLM-judge evaluations align with the empirical distribution of human grades, and are their deviations comparable to inter-rater variability?

2 Subsumption-based legal reasoning

Subsumption is the core reasoning paradigm of German civil-law doctrine (Larenz and Canaris, 1995; Möllers, 2023). Unlike common-law systems, where reasoning proceeds primarily through analogy to precedent, German legal analysis derives outcomes from codified statutory norms through a structured four-step scaffold: *Obersatz* (the norm hypothesis to be examined), *Definition*

(interpretation of the norm’s elements), *Subsumtion* (anchoring the concrete facts under those elements)², and *Ergebnis* (the resulting legal consequence). The scaffold is the German articulation of the legal syllogism that underlies statutory reasoning across civil-law jurisdictions; what is Germany-specific is its rubric-graded written form, not the reasoning operation itself. This is the *Gutachtenstil*, the opinion style in which the conclusion follows the reasoning, distinct from the result-first *Urteilssstil* used in judicial decisions; German legal education, state examinations, and entry-level practice are all trained and assessed in this register. Grading targets the written *Falllösung* (case solution) end-to-end rather than the final answer alone, because the rubric rewards how each norm element is anchored in the specific facts.

This makes subsumption a structurally demanding generation task: a system must identify the governing norm, articulate its elements, anchor every element in the facts, and conclude, a chain in which surface fluency can mask a missing *Subsumtion* step. The same structure, however, makes the task unusually amenable to rubric-based evaluation, because each step in the scaffold maps to an identifiable part of the answer rather than to a global quality impression. English-language legal NLP benchmarks built on classification, multiple-choice, or

²*Subsumtion* carries a broad and a narrow sense. Broadly, it names the entire norm-application method, the sense in the paper title. Narrowly, it names the single fact-anchoring step of the scaffold, the sense used for the scaffold step and the rubric dimension of the same name.

precedent-style reasoning (Chalkidis et al., 2022; Guha et al., 2023) do not exercise this scaffold. German terms are glossed at first use throughout; a consolidated glossary is in Appendix Q.

3 Related work

Legal NLP benchmarking has moved from classification baselines, LexGLUE (Chalkidis et al., 2022) (seven English legal datasets) and Legal-Bench (Guha et al., 2023) (162 mostly short-form US-federal-law tasks), to reasoning-centric, jurisdiction-specific evaluation. Recent work targets jurisdiction-specific reasoning: Chinese law and legal practice (Fei et al., 2024; Li et al., 2024; Shi et al., 2026), Korean bar-exam and provision-grounded reasoning (Kim et al., 2024; Lee et al., 2025), free-text Greek bar exams (Chlapanis et al., 2025), ECHR argument chains (Chlapanis et al., 2024), Thai financial and tax law (Akarajardwong et al., 2025), and US labor-law statutory simplification (Hariri and Ho, 2025), as well as retrieval-augmented evaluation (Pipitone and Houir Alami, 2024; Butler and Butler, 2026) and discrepancy-based auditing of LLM legal reasoning (Choudhury et al., 2026). The closest precedent is LEXam (Fan et al., 2026): 7,537 questions across 340 Swiss law exams in English and German, with an ensemble LLM-as-a-Judge validated against human experts, but its rubric is not the *Obersatz-Definition-Subsumtion-Ergebnis* scaffold central to German *Falllösung* grading (Section 2) and it does not couple its free-text items to the dimension-level grading criteria used in German law school and state-exam evaluation.

Free-text legal evaluation increasingly relies on LLM-as-a-Judge approaches, which raises concerns over bias, calibration, and alignment with expert judgement. German legal education shows how variable human grading already is: when Hufeld (Hufeld, 2024) had 23 graders score 15 beginner-level law exams, the average range between the lowest and highest grade per answer was 6.47 points on the 0-18 scale, and only 42% of grades fell within ± 1 point of the answer-specific mean. Legal grading is therefore better treated as a distribution over plausible expert judgements than as a single ground truth, which motivates hybrid methods that combine human expertise with calibrated LLM-based scoring.

Resources for German legal NLP remain limited: GerLayQA (Büttner and Habernal, 2024) pro-

vides ~21,000 lay questions paired with lawyer answers but focuses on QA rather than subsumption, GerLeRB (Weber et al., 2025) targets legislative retrieval, and our prior BenGER platform (Nagl and Grabmair, 2026) provides an end-to-end framework for German legal case analysis and subsumption-based reasoning workflows, on which the present dataset was built. Overall, German-law-style reasoning grounded in doctrinal subsumption under statutory law remains underserved, especially for open-ended, structured analysis under noisy human judgement. The BenGER dataset addresses this gap by combining a 596-task exam-style corpus and 531 short doctrinal reasoning items with a controlled Benchathon validation subset, a German doctrinal rubric, and a multi-rater cross-validation protocol.

4 Dataset

The BenGER dataset draws on three components covering the main forms of legal reasoning in German legal education: two exam-style case corpora, in which a fact pattern is given and a structured *Falllösung* is expected, and one set of short doctrinal reasoning items probing core principles without the full case-analysis scaffolding.

Only the first component builds on pre-existing material. The ZJS corpus contributes 581 publicly available exam-style cases from the *Zeitschrift für das Juristische Studium* (Journal for the Study of Law)³, covering civil, criminal, and public law and spanning early undergraduate exercises through first and selected second state-examination materials, each paired with an expert-written reference solution. The cases and reference solutions pre-exist this work as journal publications; our contribution is the benchmark artifact: selection, machine-readable task-reference pairs, per-item IP clearance (§Ethical considerations). The Benchathon corpus adds 15 newly collected exam-style tasks at an intermediate difficulty level (more demanding than introductory exercises, less complex than full first-state-exam problems). The tasks and their reference solutions were authored for this benchmark by the domain experts on the author team (the task creators of Appendix F). It comprises 220 human solutions in total (65 traditional, 155 human-AI co-creation; about 15 per task).⁴ The Doctrinal Principles corpus (*Grund-*

³<https://www.zjs-online.com/>

⁴The corpus is named after the in-person Benchathon

prinzipien) contributes 531 short doctrinal reasoning items (question-answer pairs with explanatory reasoning) on core principles, with recent case law where appropriate. The items were newly written for this benchmark by a subset of the authors with the relevant domain expertise and, like Benchathon, were unpublished prior to this work. The three play distinct roles: ZJS supports at-scale automated evaluation; Benchathon supports the human baseline and judge-validation layer; Doctrinal Principles probes the same reasoning at shorter form.

The benchmark is for evaluation only; we do not define a train/test split, since the objective is comparative evaluation of deployed LLM systems and human baselines. Pretraining-contamination considerations across the three corpora are discussed in §Limitations.

4.1 Benchathon task design and human solution collection

Benchathon participants solved tasks in a controlled annotation environment, with each task randomly assigned to one of two conditions. In the *traditional* condition, participants could use statutory texts, legal databases, literature, and internet research. In the *co-creation* condition, they could additionally use any LLM(s) of their choosing (commercial or open-weight, in any combination) through their own accounts. We did not log which tool a participant used. The co-creation condition therefore measures an upper bound on what human-LLM teams currently achieve with off-the-shelf assistance, not a controlled comparison against any specific system. Only the final written solution is evaluated, mirroring legal examination practice. Each task had a two-hour time limit; participant expertise was measured objectively (reported legal grades and qualifications) and subjectively (self-assessed competence on a Likert scale).

5 Experimental setup

5.1 Generation

We evaluate 12 contemporary LLM systems: 6 closed-weight through provider APIs (OpenAI, Anthropic, Google) and 6 open-weight through

event at which participants produced these solutions and the validation human grading. Participants were predominantly law students at various stages of legal training, plus a smaller group of *Referendare* (post-state-exam legal trainees), recent graduates, and a small layperson cohort; full composition in Appendix E.

a managed inference provider (DeepInfra). Per-system providers, snapshots, observed temperatures, token-budget settings, truncation counts, and the short aliases used in the body are in Table 10 (Appendix B). Each task received one generation per system in the final run. All systems generated the Benchathon subset. On ZJS, all 12 systems covered up to 581 published exams.⁵

All systems received the same German-language system and instruction prompts (Appendix J + K). Each returned a single JSON object containing the full case solution (*Falllösung*). The instruction prompt mirrors the *Gutachtenstil* scaffold (Section 2). We used provider-recommended temperatures and token budgets (temperature 0.6-1.0, max output 8,000-16,000 tokens; ZJS ran at 8,000). Generations that were malformed or shorter than 200 output tokens after a bounded retry budget were excluded. Of the ZJS generations truncated at that budget, 1,072 were re-generated at 20,000 tokens; only the replacement enters the analyses. An earlier partial Benchathon round (April 2026; different prompt scaffold and settings, partially different model set) was likewise superseded by the final run; its 28 generations from retained systems are excluded from leaderboard aggregates, and the one blind-graded validation solution drawn from it stays in the judge-validation layer (Section 5.5). Truncation in kept generations falls mostly on the smaller open-weight systems, with a few on Gemini-3.1-Pro (Table 10).

5.2 Black-box system evaluation

We evaluate LLM systems through provider APIs in a black-box setting. Our results therefore measure end-to-end system behaviour, not isolated base-model capability. We do not enable retrieval or tool-use endpoints. Components we cannot disable operate by default, notably content filters and, on some products, internal routing between sub-models. This reflects realistic deployment conditions. But it limits causal claims about underlying architectures; for open weights we test the reasoning mechanism directly with a same-base-model Thinking-vs.-Instruct toggle (Appendix P).

⁵Individual questions were occasionally flagged as prohibited content by the provider’s safety filter, leaving a small handful of per-system coverage gaps below the corpus total.

5.3 Automatic metrics

Per generation we compute BLEU (Papineni et al., 2002), ROUGE (Lin, 2004), METEOR (Banerjee and Lavie, 2005), chrF (Popović, 2015), BERTScore (Zhang et al., 2020), MoverScore (Zhao et al., 2019), and a Sentence-BERT cosine similarity (Reimers and Gurevych, 2019) against the expert reference solution. Per-system means appear in the main leaderboard tables. Per-generation values are correlated against the LLM-judge rubric score in RQ3.

5.4 LLM-as-a-Judge evaluation

For rubric-based evaluation on ZJS and Benchathon, we use a specialised LLM-as-a-Judge prompt aligned with German legal grading practice. The judge receives the task, the reference solution, and the solution to be evaluated. It scores ten dimensions: result correctness (*Ergebnisrichtigkeit*), issue identification (*Vollständigkeit*), legal grounding (*Rechtsgrundlagen*), doctrinal knowledge (*Rechtskenntnis*), subsumption quality (*Subsumtion*), problem depth (*Schwerpunktsetzung*), methodological structure (*Methodischer Stil*), organization (*Gliederung*), terminology (*Sprache*), and formal correctness (*Formalia*). Dimension scores aggregate to a 100-point raw score, mapped to the German 0-18 grade-point scale. A generation *passes* when its grade-point score reaches *ausreichend* (≥ 4 on the 0-18 scale, a raw score of $\geq 50/100$), following the German state-examination (*Staatsexamen*) convention.⁶ This is the boolean in the *Pass* columns of Table 1 and in the pass-rate aggregates throughout. The rubric operationalises the *Gutachtenstil* scaffold (Section 2): each step in the *Obersatz-Definition-Subsumtion-Ergebnis* schema is reflected in a dedicated dimension (mapping and percentage weights in Table 8, Appendix A); scaffold-step dimensions carry 90 of the 100 raw points, and a leaderboard on those seven dimensions alone reproduces the full-rubric system ranking exactly on both *Falllösung* corpora (Spearman $\rho = 1.00/1.00$; Table 9). The per-dimension scores in Figure 5 therefore track identifiable components of the reasoning process. Grundprinzipien, the short-form Doctrinal Prin-

⁶The ≥ 4 pass mark is statutory: the federal grading ordinance (Bundesministerium der Justiz, 1981) fixes *ausreichend* (4-6 of 18) as the lowest passing grade on a deliberately severe scale ($\approx 30\%$ fail the first state examination). The raw-to-grade conversion is non-linear: grade 4 requires $\geq 50/100$ raw points (table in Appendix L).

ciples corpus, uses a smaller 4-dimension rubric (result correctness, legal knowledge, subsumption, clarity). It is not mapped to the grade-point scale, since it does not follow the full *Gutachtenstil* scaffold.

The judge is blind to solution origin and receives inputs in a fixed (task, reference, candidate) order. The primary judge is GPT-5.4-mini⁷. Judge-based evaluation may itself be model-dependent. We therefore cross-validate the judge against a multi-rater human-grading layer on the Benchathon subset. We also run two stability checks on that subset: *within-judge stochasticity* ($k=3$ re-runs of GPT-5.4-mini, reported in Appendix N) and *between-judge bias* (five additional judges, listed in Section 5.5, reported in RQ5). Outputs conform to a predefined JSON schema. Non-conforming outputs are retried, and reported as evaluation failures if they cannot be parsed.

5.5 Human evaluation

Human evaluation runs on a controlled subset of 45 Benchathon solutions: one traditional human, one human-AI co-created, and one LLM-generated solution per Benchathon task. Human submissions are the median-expertise participant per legal area (*Bereich*). LLM submissions are drawn through a tier-balanced rotation across closed-flagship, efficiency-oriented, and open-weight reference systems. Each selected solution receives three independent blind reviews and one un-blind *creator review* by the task’s creator (180 reviews total). Reviewers grade against the same rubric as the LLM judge and form the IRR pool used to validate it. Creator reviews are excluded from IRR, since creators have privileged knowledge of the intended solution. The selection procedure, the seven-grader roster, and the $m = 5$ blind subset (the graders who performed blind reviews; each solution is reviewed by three of them) are described in Appendix F.

The central question is not whether the LLM judge reproduces a single human grade, but whether its deviation is comparable to that among human graders. Alongside the Calderon alt-test (Calderon et al., 2025), RQ5 reports descriptive agreement statistics: Spearman, Pearson, MAE,

⁷Configuration in Appendix M, prompt in Appendix L. An earlier version of this work used GPT-5-mini as primary judge; the final version re-scored every generation with GPT-5.4-mini and extended cross-validation to six judges, with tier-level conclusions unchanged (details in Appendix C).

Table 1: Main leaderboard across the three corpora. Per-system mean LLM-judge raw score (0-100, mean $\pm$ half-width of a 95% CI: percentile bootstrap $B = 2000$ on Benchathon for small per-system n , analytic $1.96 \cdot \text{SE}$ on ZJS and Doctrinal Principles) and pass rate. Doctrinal Principles also reports yes/no (*Ja/Nein*) decision accuracy. Systems are ranked by the mean of their available per-corpus raw scores. Per-system n sits below the corpus total where the provider’s safety filter refused individual items (§5.1). The two human-baseline rows use the LLM judge’s grades of human-written Benchathon solutions.

System / Group	Benchathon			ZJS			Doctrinal Principles			
	n	Raw	Pass	n	Raw	Pass	n	Raw	Pass	Acc.
Gemini-3.1-Pro	15	68.3 $\pm$ 8.3	93%	579	61.8 $\pm$ 1.0	83%	531	83.2 $\pm$ 1.8	87%	83%
Opus-4.7	15	69.3 $\pm$ 5.5	93%	581	58.7 $\pm$ 1.1	76%	531	83.1 $\pm$ 2.0	86%	81%
GPT-5.4	15	66.6 $\pm$ 4.9	100%	581	60.4 $\pm$ 1.1	78%	531	80.3 $\pm$ 2.1	82%	79%
Sonnet-4.6	15	58.9 $\pm$ 5.9	73%	581	51.5 $\pm$ 1.0	50%	531	81.6 $\pm$ 2.1	83%	81%
Gemini-3.1-Flash-Lite	15	63.5 $\pm$ 5.7	87%	579	44.6 $\pm$ 0.8	26%	531	73.3 $\pm$ 2.3	76%	76%
GPT-5.4-mini	15	58.6 $\pm$ 5.6	73%	581	49.7 $\pm$ 0.9	46%	531	71.8 $\pm$ 2.4	72%	71%
DeepSeek-V4-Pro	15	56.1 $\pm$ 6.1	67%	581	50.1 $\pm$ 1.0	45%	531	73.5 $\pm$ 2.3	73%	73%
DeepSeek-V4-Flash	15	49.3 $\pm$ 3.7	47%	581	45.9 $\pm$ 0.8	30%	530	70.7 $\pm$ 2.4	73%	69%
Qwen3.5-122B	15	49.4 $\pm$ 4.9	40%	581	44.2 $\pm$ 0.8	25%	526	69.8 $\pm$ 2.4	75%	78%
Qwen3.6-35B	15	42.4 $\pm$ 7.6	27%	581	39.8 $\pm$ 0.8	13%	531	69.5 $\pm$ 2.5	69%	71%
Qwen3-235B	15	41.9 $\pm$ 4.3	20%	581	35.9 $\pm$ 0.7	6%	531	67.3 $\pm$ 2.5	69%	73%
Llama-4	15	37.8 $\pm$ 4.3	20%	581	36.1 $\pm$ 0.6	4%	531	70.1 $\pm$ 2.3	75%	77%
Human (all)	220	61.1 $\pm$ 2.1	77%	—	—	—	—	—	—	—
traditional	65	50.0 $\pm$ 3.6	52%	—	—	—	—	—	—	—
co-creation	155	65.7 $\pm$ 2.1	88%	—	—	—	—	—	—	—

Cohen’s κ (Cohen, 1960), ICC(2,1) and ICC(2,k) (Shrout and Fleiss, 1979; Koo and Li, 2016). Each answers a distinct question: rank vs. absolute agreement, error magnitude, the pass/fail threshold decision, and the reliability ceiling of the human pool itself (per-statistic rationale in Appendix I). The Calderon alt-test is the primary criterion: it is purpose-built to decide whether the judge can stand in for a blind reviewer. For the alt-test we follow the blind-pool procedure at $\varepsilon = 0.15$ (Calderon’s skilled-annotator tier), with Benjamini-Yekutieli FDR at $q = 0.05$ across the m blind reviewers. We also report the single-expert variant against the un-blind creator grade, at $\varepsilon = 0.20$ because Calderon recommends the stricter expert-tier penalty when the reference grader is an expert, which the creator is. Following Calderon’s single-expert case, this variant treats the un-blind creator grade as the sole reference: each blind reviewer and the judge are scored by absolute distance to that single creator grade on every shared solution, and ω is the fraction of the m blind reviewers the judge matches or beats. So it needs one creator grade per solution, not a pool of expert votes. The judge passes when the winning rate $\omega \geq 0.5$. The primary judge is GPT-5.4-mini; we additionally re-score the same solutions with Opus-4.7, Gemini-3.1-Pro, DeepSeek-V4-Pro, Qwen3.5-397B-A17B, and Sonnet-4.6 (per-judge alt-test in Tables 23 and 24).

6 Results

Unless otherwise noted, scores in this section are the LLM judge’s rubric scores, as 0-100 raw points or 0-18 German grade points. The human-grading layer covers the 45-solution validation design: one creator review and three blind reviews per solution (180 reviews total). The released analysis pipeline regenerates every number reported below from the data files published with the paper.⁸

6.1 RQ1: Overall performance of LLM systems and humans

Table 1 gives the unified per-system view across the three corpora; per-corpus full leaderboards with surface-metric columns are in Appendix A. On Benchathon, the top of the table is a closed-flagship cluster (Opus-4.7, Gemini-3.1-Pro, and GPT-5.4 in the upper 60s on the 0-100 judge scale, with pass rates above 90%), followed by the efficiency-oriented systems in the upper 50s to lower 60s. The strongest open-weight system, DeepSeek-V4-Pro, trails the top of the flagship cluster by roughly 13 raw points. The open-weight tier spans a wide range; the nominal flagship Qwen3-235B sits 27 raw points behind the lead and underperforms its size here, despite no truncation. The same ordering broadly reproduces on ZJS, where Gemini-3.1-Pro edges ahead of GPT-

⁸<https://doi.org/10.5281/zenodo.20489788>

5.4 and the weakest open-weight systems trail the field; per-system scores are reported in Table 1.

Within-task variance is substantial. Under the GPT-5.4-mini judge, per-system standard deviations across the Benchathon tasks span roughly 8-18 raw points across all tiers. So single-task performance is a noisy signal, even for top-tier systems. On Benchathon, the human cohort scores in the mid-range of the LLM distribution under both conditions, shifted upward under co-creation (see human rows in Table 1 and Figure 2 in RQ4).

The same flagship cluster leads on the short-form Doctrinal Principles corpus, with Sonnet-4.6 joining it. Doctrinal Principles items carry a yes/no (*Ja/Nein*) decision label; across systems, decision accuracy ranges from 69% to 83%.⁹

6.2 RQ2: Scaling and system effects

Closed flagship systems from OpenAI, Anthropic, and Google occupy the top of every per-corpus leaderboard. Efficiency-tier closed systems and the strongest open-weight system (DeepSeek-V4-Pro) form an overlapping middle group. Closed-API dominance narrows below the flagship tier. By tier, the flagship-vs-open-weight gap is 20 raw points on Benchathon and narrows to 16 on ZJS. The closed-vs-open-weights gap shrinks similarly, from 18 to 12 raw points. The same direction holds on Doctrinal Principles. The flagship-vs-open-weight ordering therefore survives all three task formats and the 4-dimension rubric variant. We frame these as system-level effects, not necessarily model-level ones. Black-box API access entangles the open-closed split with provider-side prompts, decoding, post-processing, and possible hidden retrieval or tool use (see [Limitations](#)).

6.3 RQ3: Evaluation validity - automatic metric correlation

Across all three corpora, lexical surface metrics correlate with the judge’s rubric raw score more strongly than embedding-based metrics do. This runs partly against the common expectation that semantic-embedding metrics dominate lexical ones on long-form legal text. On Benchathon the lexical metrics lead: METEOR ($r=+0.61$), ROUGE ($r=+0.69$), and BLEU ($r=+0.58$). MoverScore ($r=+0.62$) is the one embedding metric on par with them. BERTScore ($r=+0.48$) and sentence-embedding similarity ($r=+0.28$) trail.

The same ordering broadly holds on ZJS and Doctrinal Principles (Appendix H: Table 15, Figure 4). This fits the highly templated structure of doctrinal German exam writing, where reference and high-quality system answers share substantial n-gram material around statutory citations and the *Obersatz-Definition-Subsumtion-Ergebnis* scaffolding.

At the per-system level, the ordering shifts. Embedding-based metrics rank systems at least as well as lexical ones. On Benchathon, ROUGE leads (Spearman +0.88), but BERTScore (+0.83) and MoverScore (+0.85) outrank BLEU (+0.69) and METEOR (+0.63). On ZJS, BERTScore is the single strongest ranker (+0.97). A leaderboard built on a strong surface or embedding metric would broadly preserve the judge’s coarse-tier ordering, but disagree on the mid-tier. This motivates a calibrated LLM judge as the primary scoring instrument.

A natural concern is that the lexical advantage could reflect a judge-specific preference for reference-like text rather than a property of the task. Two analyses on the 45-solution validation subset say otherwise. First, the lexical metrics correlate strongly with the mean blind-human grade as well (METEOR: $r\approx+0.72$ pooled). Second, we correlate the judge’s deviation from that human grade (judge – mean blind grade) with each lexical metric; the result is *negative* for all three: relative to blind human graders, the judge *under-rewards* high-overlap solutions, the opposite of a reference-style bias (Table 16 and Appendix H, with per-metric values, length-controlled partials, and small- n caveats).

6.4 RQ4: Human-AI co-creation under LLM-judge scoring

Co-creation solutions score higher than traditional ones on Benchathon, by +15.7 raw points on average, about +2.8 grade points on the 0-18 scale (Figure 2; participant-clustered 95% CI [+9.6, +22.4]). The gap holds within the 23 participants who submitted solutions in both conditions, traditional and co-creation. It repeats across all three legal areas. It is in fact a lower bound on the per-participant gain: less-experienced participants skipped traditional items more often than co-creation ones, leaving the traditional pool expertise-biased upward (full robustness arc in Appendix G). With that lift, the AI-assisted humans sit alongside the best standalone LLMs we evaluated. A TOST equiv-

⁹The corpus’s gold labels are 52% yes, 48% no.




Figure 2: Mean LLM-judge raw score (0-100) per group on the Benchathon validation subset, with 95% bootstrap CIs on each group’s mean. The two AI bars pool per-generation scores across all evaluated LLM systems (*pooled*) and across the closed-flagship tier only (*flagship*).

alence test rules out any difference of more than ± 5 raw points against the closed-flagship tier (Table 1, Appendix G). Re-graded by blind human reviewers on the RQ5 validation subset, the gap shrinks to about 14 raw points. That is +2 points smaller than the LLM-judge estimate, a small absolute discrepancy. It is consistent with the GPT-5.4-mini judge’s near-zero content-dependent calibration offset, and within the blind-pool’s inter-rater spread (Limitations).

6.5 RQ5: Grading reliability

The reliability question is simple: can the LLM judge stand in for a blind human reviewer? We answer it in three stages: how closely the judge agrees with humans per solution, whether it clears the formal substitution bar, and what its disagreements look like when they happen.

Across the validation subset, the primary GPT-5.4-mini judge agrees with the per-solution mean of the blind reviewers at Pearson $r=0.76$, pass/fail Cohen’s $\kappa=0.60$, and MAE 10.8 raw / 2.7 grade points. Averaged over solutions, the judge sits just -0.01 raw points from the blind-reviewer mean. The reviewers themselves disagree by an average within-annotation max-min spread of 23.9 raw points, comparable to the several-grade-point range Hufeld (Hufeld, 2024) documents for German legal grading. The pool reaches ICC(2,k)=0.84. So the judge sits inside the cloud of human disagreement, not outside it (Table 17).

Per-solution agreement does not by itself show that the judge can *replace* a randomly drawn re-

viewer. The Calderon alt-test sets a stricter bar: $\varepsilon = 0.15$ for the blind grader pool (*skilled*) and $\varepsilon = 0.2$ for the single un-blind creator reference (*expert*), with Benjamini-Yekutieli FDR (§5.5). We apply it across six judges: the primary GPT-5.4-mini and five additional judges, Opus-4.7, Gemini-3.1-Pro, DeepSeek-V4-Pro, Qwen3.5-397B-A17B, and Sonnet-4.6.

For the blind pool, two judges clear the $\omega \geq 0.5$ replacement bar on the human-written solutions: GPT-5.4-mini at $\omega = 0.60$ and Opus-4.7 at $\omega = 0.60$. Sonnet-4.6 ($\omega = 0.40$) and Gemini-3.1-Pro ($\omega = 0.40$) fall one annotator-rejection short. DeepSeek-V4-Pro and Qwen3.5-397B-A17B fail at $\omega \leq 0.20$. At $m = 5$ these are point estimates; the two-judge pass shows a direction, but the reviewer sample is too small to establish it statistically (Limitations).

Against the single-expert creator reference, GPT-5.4-mini passes on human-written solutions ($\omega = 0.60$) but fails on LLM-generated ones ($\omega = 0.00$) (Table 17; per-judge ω , per-annotator p -values, and a bootstrap robustness check in Appendix N).

Agreement is highest on scoring dimensions with structured criteria and lowest on presentation and form (Figure 5). Several outcome and reasoning dimensions drop on LLM-generated solutions. Result correctness goes from $r=0.76$ (human-written) to $r=0.62$ (LLM), with comparable or larger drops on Issue identification, Doctrinal knowledge, and Problem depth.

6.6 Error analysis

On Benchathon and ZJS, the two free-text corpora, we find two dominant failure modes, one typical of the flagship tier and one of the open-weight tier. (i) *Long-but-shallow generations*, concentrated in the flagship tier, run to over 2,000 words yet score at most half marks on *Subsumtion* and problem depth. The pattern is over-writing, which the rubric does not reward. (ii) *Wrong-norm framing generations*, concentrated in open-weight systems, are methodologically well-formed but score low on legal grounding and doctrinal knowledge. These solutions argue from the wrong legal norms rather than making mistakes while applying the right ones. The two modes are largely distinct (119 of the ~1,800 flagged ZJS generations fall in both). Counts of each failure mode, and how often it occurs per tier, are in Appendix O for all three corpora.

On the short-form Doctrinal Principles corpus, yes/no accuracy follows the same tier ordering, and the wrong decision itself is the dominant error. Fact-bound *Subsumtion* and outcome correctness are tightly coupled on every corpus, so reasoning-outcome divergence is rare but does occur. Where it occurs, the dimension scores show a consistency failure rather than a reasoning failure: the model reaches a wrong decision despite sound *Subsumtion*, or the right decision on flawed *Subsumtion*. Its stated answer is decoupled from its own reasoning, an inconsistency that yes/no accuracy cannot capture but the rubric does.

7 Discussion

Across the three corpora, Opus-4.7, Gemini-3.1-Pro, and GPT-5.4 consistently form the top of the leaderboard, with Sonnet-4.6 joining them on Doctrinal Principles. The flagship-vs-open-weight gap is ~20 raw points on Benchathon (RQ2), narrower on the longer ZJS corpus, and most compressed on Doctrinal Principles. The strongest open-weight system, DeepSeek-V4-Pro, scores within the range of the closed efficiency-oriented systems on the longer corpora and comes close to the flagships on the short one, so the practical question for German legal NLP is not whether open-weight systems can be used, but whether the residual flagship advantage justifies the open- versus closed-weight access constraints (Appendix B).

Surface metrics (BLEU, ROUGE, METEOR) correlate with the rubric judge more strongly than embedding metrics do (RQ3), consistent with the templated *Obersatz-Definition-Subsumtion-Ergebnis* structure of German exam writing. A lexical-only leaderboard would keep the coarse tiers but disagree on the mid-tier, so the rubric judge remains the primary instrument and surface metrics read as cheap correlates.

Human-AI co-creation lifts unaided humans into the closed-flagship score range, with a magnitude caveat: AI-assisted human work scores +15.7 raw points ($\approx +2.8$ German grade points) above unaided human work under the judge (RQ4), the gap holds within-participant and across the three legal areas, and a TOST equivalence test renders it statistically indistinguishable from the closed-flagship LLM tier (Appendix G). Re-graded by blind human reviewers, the gap comes out about 2 raw points smaller, with the same sign and the same conclusion. That difference is far smaller than

the disagreement among the human graders themselves.

The primary GPT-5.4-mini judge tracks the per-solution blind-reviewer mean at Pearson $r=0.76$ and sits inside the human disagreement cloud (RQ5). On the Calderon alt-test at $\epsilon = 0.15$ it also clears the $\omega \geq 0.5$ replacement bar on human-written solutions ($\omega = 0.60$), as does Opus-4.7 ($\omega = 0.60$), while the other four judges fall short; at $m = 5$ reviewers this is a point estimate, not yet statistically established (see [Limitations](#)). Rank agreement across the six judges is high: re-ranking the 12 systems under each cross-validation judge yields system-level Spearman ρ of 0.92-0.97 against the primary judge (Table 20), so the RQ1-RQ2 ordering survives any single judge’s score bias. That the two passing judges come from different providers, and agree on which blind reviewers they match, suggests same-provider bias is not a large confound. Re-scoring the leaderboard with Opus-4.7 as primary judge is a natural future robustness check.

LLM judges are thus scalable raters alongside humans rather than standalone ground truth. Where the judge still diverges, the disagreement concentrates on presentation and on outcomes reached via reasoning paths other than the reference solution’s: agreement is highest on the outcome and issue-identification dimensions, lowest on formal correctness and terminology, and slips further on LLM-generated solutions.

8 Conclusion

We introduce the BenGER dataset for subsumption-based legal reasoning in German law. It combines two exam-style corpora (ZJS, Benchathon) with a short doctrinal-reasoning corpus (Doctrinal Principles), a Benchathon validation subset with human and human-AI baselines, a rubric-aligned LLM-as-a-Judge, and a multi-rater cross-validation protocol. It comprises 12 systems and 13514 generations across the three corpora, all run end-to-end on our BenGER platform (Nagl and Grabmair, 2026). BenGER is released as an open evaluation framework with corpora, judge prompts, rubric, and analysis pipeline. The methodology transfers to any expert-graded free-text domain: treat grading as a distribution over expert assessments, validate a judge by reviewer substitution, and swap in the domain’s own rubric.

Limitations

All LLM systems are evaluated through hosted provider APIs in a black-box setting, so reported scores reflect end-to-end deployed behaviour, not isolated base-model capability. We do not vary provider-side prompts, decoding configurations, retrieval, or tool use, and make no claim about the contribution of individual architectural choices.

The dimension rubric and the LLM-judge prompt were co-developed for this benchmark, not adopted from independent prior work. We validate them as a pair against the multi-rater human-grading layer (RQ5). Its within-annotation inter-rater spread (5.2 grade points) sits within the range Hufeld (Hufeld, 2024) documents for German legal grading. Neither component has been validated in isolation.

The Benchathon subset comprises 15 tasks and a modest participant pool. This suffices for the per-task, per-system comparisons and the judge validation reported here. But it limits statistical power for fine-grained interaction effects, such as expertise-stratified co-creation analyses.

Judge bias is a known confounder in LLM evaluation (Zheng et al., 2023; Panickssery et al., 2024). We address it via the multi-rater agreement protocol on the Benchathon validation subset and the inter-judge agreement analysis (RQ5): the six judges differ in how strictly they score, but their pairwise rank agreement is comparable to the agreement among the human reviewers.

Same-provider preference might occur, but its plausible size is limited, for four reasons. First, re-ranking the systems under every cross-validation judge preserves the ordering across providers (system-level $\rho \geq 0.92$, Table 20). Second, judge-human agreement is *lower* on the LLM-generated half of the validation subset (Pearson $r=0.64$ vs. $r=0.76$), the opposite of what same-family inflation would produce. Third, the Calderon winning rate drops from $\omega = 0.60$ on human-written solutions to 0.00 on LLM-generated ones, the same direction. Fourth, that drop appears for all six judges tested, closed and open-weight, including judges whose model families lead the system rankings, so the pattern is broad across judge models rather than aligned with any one family (Tables 21, 23, 24).

Our per-annotator sample sizes ($n_j \in [7, 21]$ scored solutions per blind reviewer) are well below the 50-100 Calderon recommends. Every test therefore fell back to Wilcoxon, with reduced

power. At this sample size the two-judge alt-test pass is a point estimate, not a robust result. The winning rate moves in steps of 0.20 at $m = 5$ reviewers. The 95% CI on the primary judge’s ω spans [0.15, 0.95]. No bootstrap lower bound clears $\omega = 0.5$ for any judge (Appendix N). The two-judge pass should therefore be read as a promising direction, not an established result.

Absolute RQ1-RQ2 leaderboard scores carry the GPT-5.4-mini calibration offset (Table 21, measured against the blind-reviewer pool mean). They should be read as a calibrated ordering, not human-pool estimates. Calderon’s FAQ recommends computing ω per domain, but splitting by legal area would leave each blind reviewer only 3-7 scored solutions per area (civil, public, and criminal law), so we pool across the three areas.

The rubric is grounded in German doctrinal grading conventions and is not designed to generalise to common-law or non-German civil-law systems. Metrics and aggregate scores in the BenGER dataset should not be transferred to other jurisdictions.

ZJS material is publicly indexed, so provider models may have encountered some of it during pretraining; we have not run an explicit memorisation probe. The Benchathon and Doctrinal Principles corpora were unpublished at the time of all generation runs. They are not a clean contamination control: their orderings agreeing with the ZJS leaderboard is consistent both with low contamination and with contamination that scales proportionally across systems.

Per-area CIs (RQ4), pairwise inter-judge correlations (RQ5), and error-analysis patterns are reported descriptively. We draw no inferential claims from individual cell-level comparisons.

The Benchathon and ZJS exam-style corpora are graded with the 10-dimension case-solution (*Falllösung*) rubric. The short-form Doctrinal Principles corpus uses a 4-dimension rubric (result correctness, legal knowledge, subsumption, clarity). Per-dimension comparisons across corpora are therefore not meaningful. Only aggregate raw scores, pass rates, and yes/no accuracy (Doctrinal Principles only) should be compared across the three. The Google efficiency-tier slot additionally mixes two Google models across corpora; Appendix B details the change and how to read cross-corpus comparisons for it.

Ethical considerations

Benchathon participants were recruited from a pool of law students and recent graduates who consented to the use of their written solutions and self-reported expertise data for research purposes. Solutions and self-reports are anonymized in the released BenGER dataset. Human grading is conducted by domain experts under the same consent regime. We release dimension-level rubric scores and aggregate statistics, not free-text feedback that could re-identify a participant. The Benchathon study design for the human and co-creation baseline was approved by an ethics committee. The LLM-as-a-Judge methodology in BenGER is intended for research-scale evaluation of legal-reasoning systems. It is not validated for use as a graded assessment instrument in legal education, or for deployment in legal practice.

The ZJS exams and reference solutions are copyright-protected. We have obtained re-publishing consent for the vast majority of items, bundled directly with the released dataset. For the remainder we provide URL pointers to publicly accessible free locations on the web. All model generations and LLM-judge evaluation scores are released under a Creative Commons Attribution 4.0 (CC BY 4.0) licence. The benchmark is therefore end-to-end reproducible: recomputing any reported number requires only a short fetch step for the ZJS items without re-publishing consent.

A public legal-reasoning benchmark can be optimised against. Future systems may improve on the benchmark without commensurate gains in real legal-reasoning quality. To mitigate this, we release the rubric and judge prompts alongside the data, encourage reporting on the full per-dimension breakdown rather than a single headline score, and treat the benchmark as one signal among several in the evaluation of legal-domain LLMs.

AI usage

The authors used Anthropic Claude Opus 4.7, OpenAI GPT-5.5 Pro, and Anthropic Claude Fable 5 (camera-ready revision) for coding, drafting, and writing assistance throughout. AI agents implemented parts of the analysis code; the authors checked the implementations and verified all reported results. Generated text was checked, edited, and approved by the authors. The authors take full responsibility for the paper’s contents.

Acknowledgments

We thank our colleagues in the TITAN project group - Philipp Reuß (University of Göttingen), Sarah Rachut (TU Braunschweig), Liane Wörner (University of Konstanz), Dominik Brodowski (Saarland University), and Markus Langer and Asya Caroei (both University of Freiburg) - for the discussions that shaped this work. We are grateful to the Legal Tech Verband (LTV) and its members Clemens Hufeld (Beck/Noxtua), Daniella Domokos (Capgemini), Ben Kühnel (Forvis Mazars), and Katharina Hertel (Recode Law) for their support in organising the Benchathon, and to all Benchathon participants for their dedicated work. We thank Markus Wagner and the *Zeitschrift für das Juristische Studium* (ZJS) author community for generously permitting use and redistribution of their cases. Finally, we thank Peter Moser (LMU) for connecting us with the right experts and materials.

BenGER is generously funded by the Daimler Benz Foundation as part of project TITAN (*Technologische Intelligenz zur Transformation, Automatisierung und Nutzerorientierung des Justizsystems*) and by the German Federal Ministry of Justice and Consumer Protection as part of the *Digitalisierungsinitiative des Bundes für die Justiz* under the GSJ project (*Generatives Sprachmodell der Justiz*), a collaboration between the justice ministries of North Rhine-Westphalia and Bavaria, TU Munich, and the University of Cologne.

References

- Pawitsapak Akarajadwong, Pirat Pothavorn, Chompakorn Chaksangchaichot, Panuthep Tasawong, Thitiwat Nopparatbundit, Keerakiat Pratai, and Sarana Nutanong. 2025. [NitiBench: Benchmarking LLM frameworks on Thai legal question answering capabilities](#). In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pages 34304–34327, Suzhou, China. Association for Computational Linguistics.
- Satanjeev Banerjee and Alon Lavie. 2005. [METEOR: An automatic metric for MT evaluation with improved correlation with human judgments](#). In *Proceedings of the ACL Workshop on Intrinsic and Extrinsic Evaluation Measures for Machine Translation and/or Summarization*, pages 65–72, Ann Arbor, Michigan. Association for Computational Linguistics.
- Bundesministerium der Justiz. 1981. [Verordnung über eine Noten- und Punkteskala für die erste und zweite juristische Prüfung \(JurPrNotSkV\)](#). BGBl. I S. 1243.

- Federal grading ordinance for the two German state law examinations.
- Abdur-Rahman Butler and Umar Butler. 2026. [Legal RAG Bench: An end-to-end benchmark for legal RAG](#). *Preprint*, arXiv:2603.01710.
- Marius Büttner and Ivan Habernal. 2024. [Answering legal questions from laymen in German civil law system](#). In *Proceedings of the 18th Conference of the European Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 2015–2027, St. Julian’s, Malta. Association for Computational Linguistics.
- Nitay Calderon, Roi Reichart, and Rotem Dror. 2025. [The alternative annotator test for LLM-as-a-judge: How to statistically justify replacing human annotators with LLMs](#). In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 16051–16081, Vienna, Austria. Association for Computational Linguistics.
- Ilias Chalkidis, Abhik Jana, Dirk Hartung, Michael Bommarito, Ion Androutsopoulos, Daniel Katz, and Nikolaos Aletras. 2022. [LexGLUE: A benchmark dataset for legal language understanding in English](#). In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 4310–4330, Dublin, Ireland. Association for Computational Linguistics.
- Odyseas S. Chlapanis, Dimitrios Galanis, Nikolaos Aletras, and Ion Androutsopoulos. 2025. [GreekBar-Bench: A challenging benchmark for free-text legal reasoning and citations](#). In *Findings of the Association for Computational Linguistics: EMNLP 2025*, pages 25099–25119, Suzhou, China. Association for Computational Linguistics.
- Odyseas S. Chlapanis, Dimitrios Galanis, and Ion Androutsopoulos. 2024. [LAR-ECHR: A new legal argument reasoning task and dataset for cases of the European Court of Human Rights](#). In *Proceedings of the Natural Legal Language Processing Workshop 2024*, pages 267–279, Miami, FL, USA. Association for Computational Linguistics.
- Manan Roy Choudhury, Adithya Chandramouli, Manan Anand, and Vivek Gupta. 2026. [Better call CLAUSE: A discrepancy benchmark for auditing LLMs legal reasoning capabilities](#). In *Findings of the Association for Computational Linguistics: EACL 2026*, pages 5776–5818, Rabat, Morocco. Association for Computational Linguistics.
- Jacob Cohen. 1960. [A coefficient of agreement for nominal scales](#). *Educational and Psychological Measurement*, 20(1):37–46.
- Yu Fan, Jingwei Ni, Jakob Merane, Yang Tian, Yoan Hermstrüwer, Yinya Huang, Mubashara Akhtar, Etienne Salimbeni, Florian Geering, Oliver Dreyer, Daniel Brunner, Markus Leippold, Mrinmaya Sachan, Alexander Stremitzer, Christoph Engel, Elliott Ash, and Joel Niklaus. 2026. [LEXam: Benchmarking legal reasoning on 340 law exams](#). In *Proceedings of the Fourteenth International Conference on Learning Representations*. ArXiv:2505.12864.
- Zhiwei Fei, Xiaoyu Shen, Dawei Zhu, Fengzhe Zhou, Zhuo Han, Alan Huang, Songyang Zhang, Kai Chen, Zhixin Yin, Zongwen Shen, Jidong Ge, and Vincent Ng. 2024. [LawBench: Benchmarking legal knowledge of large language models](#). In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 7933–7962, Miami, Florida, USA. Association for Computational Linguistics.
- Neel Guha, Julian Nyarko, Daniel E. Ho, Christopher Ré, Adam Chilton, Aditya Narayana, Alex Chohlas-Wood, Austin Peters, Brandon Waldon, Daniel N. Rockmore, Diego Zambrano, Dmitry Talisman, Enam Hoque, Faiz Surani, Frank Fagan, Galit Sarfaty, Gregory M. Dickinson, Haggai Porat, Jason Hegland, and 21 others. 2023. [LegalBench: A collaboratively built benchmark for measuring legal reasoning in large language models](#). In *Advances in Neural Information Processing Systems*, volume 36, pages 44123–44279. Curran Associates, Inc.
- Emaan Hariri and Daniel E. Ho. 2025. [AI for statutory simplification: A comprehensive state legal corpus and labor benchmark](#). In *Proceedings of the Twentieth International Conference on Artificial Intelligence and Law*, pages 177–187, Chicago, IL, USA. ACM. ArXiv:2508.19365.
- Clemens Hufeld. 2024. [Jede Korrektur eine andere Note: Quantitative Untersuchung der Objektivität juristischer Klausurbewertungen](#). *Zeitschrift für Diaktik der Rechtswissenschaft*, 11(1):59–83.
- Zheng Jia, Shengbin Yue, Wei Chen, Siyuan Wang, Yidong Liu, Zejun Li, Yun Song, and Zhongyu Wei. 2026. [Ready jurist one: Benchmarking language agents for legal intelligence in dynamic environments](#). In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 10351–10376. Association for Computational Linguistics.
- Yeeun Kim, Youngrok Choi, Eunkyoung Choi, JinHwan Choi, Hai Jin Park, and Wonseok Hwang. 2024. [Developing a pragmatic benchmark for assessing Korean legal language understanding in large language models](#). In *Findings of the Association for Computational Linguistics: EMNLP 2024*, pages 5573–5595, Miami, Florida, USA. Association for Computational Linguistics.
- Terry K. Koo and Mae Y. Li. 2016. [A guideline of selecting and reporting intraclass correlation coefficients for reliability research](#). *Journal of Chiropractic Medicine*, 15(2):155–163.
- Karl Larenz and Claus-Wilhelm Canaris. 1995. [Methodenlehre der Rechtswissenschaft](#), 3 edition. Springer, Berlin, Heidelberg.

- Jihyung Lee, Daehui Kim, Seonjeong Hwang, Hyounghun Kim, and Gary Lee. 2025. [KoBLEX: Open legal question answering with multi-hop reasoning](#). In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pages 4019–4053, Suzhou, China. Association for Computational Linguistics.
- Haitao Li, You Chen, Qingyao Ai, Yueyue Wu, Ruizhe Zhang, and Yiqun Liu. 2024. [LexEval: A comprehensive Chinese legal benchmark for evaluating large language models](#). In *Advances in Neural Information Processing Systems*, volume 37, pages 25061–25094. Curran Associates, Inc.
- Chin-Yew Lin. 2004. [ROUGE: A package for automatic evaluation of summaries](#). In *Text Summarization Branches Out*, pages 74–81, Barcelona, Spain. Association for Computational Linguistics.
- Thomas M. J. Möllers. 2023. *Juristische Methodenlehre*, 5 edition. C.H. Beck, München.
- Sebastian Nagl and Matthias Grabmair. 2026. BenGER Platform: A collaborative web platform for end-to-end benchmarking of German legal tasks. In *Proceedings of the Twenty-First International Conference on Artificial Intelligence and Law*, Singapore. ACM. To appear. arXiv:2604.13583.
- Arjun Panickssery, Samuel R. Bowman, and Shi Feng. 2024. [LLM evaluators recognize and favor their own generations](#). In *Advances in Neural Information Processing Systems*, volume 37, pages 68772–68802. Curran Associates, Inc.
- Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. 2002. [Bleu: A method for automatic evaluation of machine translation](#). In *Proceedings of the 40th Annual Meeting of the Association for Computational Linguistics*, pages 311–318. Association for Computational Linguistics.
- Nicholas Pipitone and Ghita Hour Alami. 2024. [LegalBench-RAG: A benchmark for retrieval-augmented generation in the legal domain](#). Preprint, arXiv:2408.10343.
- Maja Popović. 2015. [chrF: Character n-gram F-score for automatic MT evaluation](#). In *Proceedings of the Tenth Workshop on Statistical Machine Translation*, pages 392–395, Lisbon, Portugal. Association for Computational Linguistics.
- Nils Reimers and Iryna Gurevych. 2019. [Sentence-BERT: Sentence embeddings using Siamese BERT-Networks](#). In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 3982–3992. Association for Computational Linguistics.
- Yuzhen Shi, Huanghai Liu, Yiran Hu, Gaojie Song, Xinran Xu, Yubo Ma, Tianyi Tang, Li Zhang, Qingjing Chen, Di Feng, Wenbo Lv, Weiheng Wu, Kexin Yang, Sen Yang, Wei Wang, Rongyao Shi, Yuanyang Qiu, Yuemeng Qi, Jingwen Zhang, and 11 others. 2026. [PLawBench: A rubric-based benchmark for evaluating LLMs in real-world legal practice](#). In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 10067–10116. Association for Computational Linguistics.
- Patrick E. Shrout and Joseph L. Fleiss. 1979. [Intraclass correlations: Uses in assessing rater reliability](#). *Psychological Bulletin*, 86(2):420–428.
- Malte Weber, Balaramakrishna Paritala, Abhilash Reddy Rechu, Leila Feddoul, Suresh Kumar Bonagiri, Norman Klewer, Pirmin Mathias Karg, Christoph Unger, Marianne Mauch, and Birgitta König-Ries. 2025. [GerLeRB – German legislative retrieval benchmark](#). In *Fachtagung Rechts- und Verwaltungsinformatik (RVI 2025)*, Recht und Informatik, pages 157–168, Bonn. Gesellschaft für Informatik.
- Tianyi Zhang, Varsha Kishore, Felix Wu, Kilian Q. Weinberger, and Yoav Artzi. 2020. [BERTScore: Evaluating text generation with BERT](#). In *International Conference on Learning Representations*. ArXiv:1904.09675.
- Wei Zhao, Maxime Peyrard, Fei Liu, Yang Gao, Christian M. Meyer, and Steffen Eger. 2019. [MoverScore: Text generation evaluating with contextualized embeddings and earth mover distance](#). In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 563–578. Association for Computational Linguistics.
- Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica. 2023. [Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena](#). In *Advances in Neural Information Processing Systems*, volume 36, pages 46595–46623. Curran Associates, Inc.

Contents of the Appendix

A	Full results table	14
B	Evaluated systems catalogue	14
C	Operational metadata and reproducibility	
	14	
D	Per-task per-system heatmap	15
E	Benchathon participant composition ...	16
F	Full Benchathon human-evaluation assignment procedure	19
G	RQ4 supplementary detail	20
H	Automatic-metric vs. LLM-judge correlation	21
I	Judge-human agreement statistics	21

J	System prompts for generation	21
K	Instruction prompts for generation	23
L	Evaluation prompts for the LLM judge	26
M	LLM-judge configuration	30
N	Judge calibration - supplementary detail	30
O	Error-analysis mode counts	31
P	Reasoning-mode ablation	31
Q	Glossary of German terms	31

A Full results table

The full per-system leaderboards are split into six tables, three corpora $\times$ two metric groups: rubric-based metrics (LLM-judge raw / grade-points / pass rate; plus human-grader rubric columns on Benchathon) in Tables 2, 3, 4, and surface / embedding-based automatic metrics (plus Doctrinal Principles yes/no decision accuracy, a task-correctness measure rather than a rubric one) in Tables 5, 6, 7. The two Benchathon human-baseline rows (traditional, co-creation) appear at the top of the Benchathon tables. The CI is the percentile interval over $B = 2000$ resamples on the per-generation values (seed 42); each cell shows the mean stacked above $\pm$ half the width of the 95% CI (in a smaller font). Empty cells indicate the metric is not computed on that dataset.

A.1 Reasoning-core robustness

Table 8 maps each rubric dimension to the *Gutachtenstil* scaffold step it grades, with its percentage weight. Table 9 recomputes each system’s mean score using only the seven scaffold-step rubric dimensions (90 of 100 points; dropping *Gliederung*, *Sprache*, *Formalia*), rescaled to /100. The full-rubric system ranking is reproduced exactly on both corpora.

B Evaluated systems catalogue

Table 10 catalogues the LLM systems evaluated in this work alongside their provider family, the short alias used throughout the body of the paper, the full model id, tier classification (closed flagship, efficiency-oriented, or open-weight reference), weights status (open vs. closed), the decoding temperature and max-output-tokens setting used, the total generation count on the Benchathon subset, and the count of truncated generations over that total.

The Google efficiency-tier slot was generated with two distinct Google models: Gemini-3-Flash

on the Benchathon subset, and Gemini-3.1-Flash-Lite (a separate efficiency product line, not a successor release of Flash) on the ZJS and Doctrinal Principles corpora. The model set was finalised between the two generation runs. We treat them as the same conceptual tier slot for tier-level comparisons and disclose the raw model ids in Table 10. Row-level cross-corpus comparisons for this slot should be read with this model change in mind, and tier-level claims should not be read as claims about either Flash or Flash-Lite alone.

C Operational metadata and reproducibility

Table 11 reports per-system operational metadata (mean per-generation cost, latency, and output length) from the per-call response `_metadata` the platform logs for every request. Cost cells are empty where the provider route did not report billing at generation time.

Reproducibility addendum. The same per-call metadata records the run settings beyond the per-system temperatures and token budgets of Table 10. *Run dates*: Benchathon: 2026-05-08 to 2026-05-08. ZJS: 2026-05-13 to 2026-05-14. Doctrinal Principles: 2026-05-19 to 2026-05-19. The ZJS range spans the initial 8,000-token run and the next-day re-generation of truncated attempts; superseded attempts (the ZJS first attempts and the April Benchathon round; §Experimental setup) are excluded from all per-generation aggregates here. *Retry budget*: generation calls run under the worker-side retry budget of 3 (the same default as the judge calls in Appendix M); observed retries across kept generation calls: 0. *Sampling parameters*: temperature and max output tokens were set explicitly per system (Table 10); no other sampling parameter (top-p, top-k, seed, thinking budget, reasoning effort) appears in any logged call; those ran at provider defaults, including the reasoning behaviour of the one explicitly reasoning-tagged system (Qwen3-235B-A22B-Thinking-2507). *Earlier version*: an earlier version of this work used GPT-5-mini as the primary judge; the final version re-scored every generation with GPT-5.4-mini and extended cross-validation to six judges, leaving tier-level conclusions unchanged while absolute scores shift as a calibrated ordering (Limitations). That version’s ZJS and Benchathon aggregates also double-counted superseded generation attempts; deduplicating raises truncation-

Table 2: Rubric-based results on the **Benchathon** corpus: per-system bootstrap 95% CIs ($B = 2000$, seed 42) on the LLM-judge raw / grade-points / pass rate columns plus the parallel human-grader rubric columns (H.-prefixed). Each cell shows the mean stacked above $\pm$ half the width of the 95% CI. n is the per-system generation count; n_H is the per-system human-grader row count. The two human-baseline rows (traditional, co-creation) sit at the top.

System	n	n_H	Judge raw	Gr. pts	Pass	H. raw	H. Gr.pts	H. pass
<i>Human (traditional)</i>	65	60	50.0 ± 3.7	4.72 ± 0.76	52.3% ± 11.5	50.7 ± 5.6	5.67 ± 1.04	63.3% ± 11.6
<i>Human (co-creation)</i>	155	60	65.7 ± 2.2	8.50 ± 0.58	87.7% ± 5.2	65.4 ± 5.7	9.08 ± 1.25	86.7% ± 8.4
Opus-4.7	15	8	69.3 ± 5.7	9.47 ± 1.80	93.3% ± 10.0	62.7 ± 7.2	7.62 ± 1.75	87.5% ± 18.8
Gemini-3.1-Pro	15	8	68.3 ± 8.4	9.60 ± 1.83	93.3% ± 10.0	76.2 ± 6.4	11.62 ± 1.88	100.0% ± 0.0
GPT-5.4	15	4	66.6 ± 4.8	8.47 ± 1.43	100.0% ± 0.0	61.1 ± 16.3	7.50 ± 4.25	75.0% ± 37.5
Gemini-3.1-Flash-Lite	15	4	63.5 ± 6.1	7.73 ± 1.80	86.7% ± 16.6	67.0 ± 14.1	9.25 ± 3.38	75.0% ± 37.5
Sonnet-4.6	15	4	58.9 ± 6.0	6.47 ± 1.66	73.3% ± 23.3	53.8 ± 6.5	5.25 ± 1.75	75.0% ± 37.5
GPT-5.4-mini	15	4	58.6 ± 5.7	6.27 ± 1.53	73.3% ± 23.3	54.6 ± 10.6	5.50 ± 2.25	75.0% ± 37.5
DeepSeek-V4-Pro	15	4	56.1 ± 5.8	5.73 ± 1.47	66.7% ± 23.4	77.6 ± 6.1	11.75 ± 1.75	100.0% ± 0.0
Qwen3.5-122B	15	4	49.4 ± 4.9	4.33 ± 1.07	40.0% ± 23.4	47.1 ± 6.9	3.75 ± 1.12	25.0% ± 37.5
DeepSeek-V4-Flash	15	4	49.3 ± 3.6	4.00 ± 0.73	46.7% ± 26.6	52.4 ± 16.8	5.75 ± 3.38	75.0% ± 37.5
Qwen3.6-35B	15	0	42.4 ± 7.2	3.60 ± 1.20	26.7% ± 20.0	—	—	—
Qwen3-235B	15	8	41.9 ± 4.6	3.00 ± 0.51	20.0% ± 20.0	25.5 ± 9.5	1.50 ± 0.69	0.0% ± 0.0
Llama-4	15	8	37.8 ± 4.4	2.60 ± 0.40	20.0% ± 20.0	19.9 ± 5.1	0.88 ± 0.50	0.0% ± 0.0

Table 3: Rubric-based results on the **ZJS** corpus: per-system bootstrap 95% CIs ($B = 2000$, seed 42) on the LLM-judge raw, grade-points, and pass-rate columns. Each cell shows the mean stacked above $\pm$ half the width of the 95% CI. n is the per-system generation count.

System	n	Judge raw	Gr. pts	Pass
Gemini-3.1-Pro	579	61.7 ± 1.0	7.33 ± 0.27	83.2% ± 3.1
GPT-5.4	581	60.4 ± 1.1	6.99 ± 0.28	78.5% ± 3.5
Opus-4.7	581	58.7 ± 1.0	6.54 ± 0.27	75.7% ± 3.5
Sonnet-4.6	581	51.6 ± 1.0	4.87 ± 0.23	50.4% ± 4.2
DeepSeek-V4-Pro	581	50.0 ± 1.0	4.53 ± 0.23	44.6% ± 4.0
GPT-5.4-mini	581	49.7 ± 1.0	4.43 ± 0.21	46.1% ± 4.1
DeepSeek-V4-Flash	581	46.0 ± 0.8	3.68 ± 0.17	30.1% ± 3.6
Gemini-3.1-Flash-Lite	579	44.6 ± 0.8	3.48 ± 0.16	26.3% ± 3.5
Qwen3.5-122B	581	44.2 ± 0.8	3.43 ± 0.15	25.6% ± 3.7
Qwen3.6-35B	581	39.8 ± 0.9	2.84 ± 0.12	13.4% ± 2.8
Llama-4	581	36.2 ± 0.6	2.36 ± 0.07	4.6% ± 1.7
Qwen3-235B	581	35.9 ± 0.7	2.41 ± 0.09	6.2% ± 1.9

prone ZJS means by up to ~3 raw points and makes Benchathon aggregates final-run-only ($n=15$).

D Per-task per-system heatmap

Figure 3 visualises the per-task per-system mean LLM-judge raw score on the Benchathon subset. Rows are systems ordered by overall mean (top = highest); columns are the 15 Benchathon tasks or-

Table 4: Rubric-based results on the **Doctrinal Principles** corpus: per-system bootstrap 95% CIs ($B = 2000$, seed 42) on the LLM-judge raw, grade-points, and pass-rate columns. Each cell shows the mean stacked above $\pm$ half the width of the 95% CI. n is the per-system generation count.

System	n	Judge raw	Gr. pts	Pass
Gemini-3.1-Pro	531	83.2 ± 1.8	—	87.0% ± 2.9
Opus-4.7	531	83.1 ± 1.9	—	85.9% ± 2.9
Sonnet-4.6	531	81.6 ± 2.0	—	82.9% ± 3.2
GPT-5.4	531	80.3 ± 2.1	—	81.5% ± 3.2
DeepSeek-V4-Pro	531	73.5 ± 2.3	—	72.7% ± 3.8
Gemini-3.1-Flash-Lite	531	73.3 ± 2.2	—	75.5% ± 3.6
GPT-5.4-mini	531	71.8 ± 2.4	—	71.6% ± 3.8
DeepSeek-V4-Flash	530	70.7 ± 2.3	—	73.4% ± 3.7
Llama-4	531	70.1 ± 2.4	—	75.0% ± 3.6
Qwen3.5-122B	526	69.8 ± 2.3	—	75.3% ± 3.6
Qwen3.6-35B	531	69.5 ± 2.5	—	68.9% ± 3.8
Qwen3-235B	531	67.3 ± 2.4	—	69.5% ± 4.0

Table 5: Automatic-metric results on the **Benchathon** corpus: per-system bootstrap 95% CIs ($B = 2000$, seed 42) on the surface and embedding-based metrics. Each cell shows the mean stacked above $\pm$ half the width of the 95% CI; dashes mark metrics not computed on this dataset.

System	n	BLEU	ROUGE	METEOR	BERTSc	MovSc	SemSim
<i>Human (traditional)</i>	65	0.055 ± 0.014 0.083	0.171 ± 0.013 0.186	0.181 ± 0.020 0.255	0.325 ± 0.029 0.338	0.908 ± 0.005 0.916	0.705 ± 0.056 0.761
<i>Human (co-creation)</i>	155	0.065 ± 0.005 0.065	0.193 ± 0.006 0.193	0.203 ± 0.012 0.203	0.384 ± 0.016 0.384	0.916 ± 0.003 0.916	0.795 ± 0.025 0.795
Opus-4.7	15	0.070 ± 0.017 0.070	0.207 ± 0.021 0.207	0.202 ± 0.027 0.202	0.399 ± 0.032 0.399	0.921 ± 0.004 0.921	0.790 ± 0.049 0.790
Gemini-3.1-Pro	15	0.101 ± 0.020 0.101	0.194 ± 0.030 0.194	0.269 ± 0.036 0.269	0.363 ± 0.040 0.363	0.920 ± 0.008 0.920	0.789 ± 0.064 0.789
GPT-5.4	15	0.055 ± 0.019 0.055	0.191 ± 0.018 0.191	0.185 ± 0.027 0.185	0.384 ± 0.046 0.384	0.920 ± 0.005 0.920	0.794 ± 0.050 0.794
Gemini-3.1-Flash-Lite	15	0.094 ± 0.015 0.094	0.180 ± 0.018 0.180	0.270 ± 0.022 0.270	0.338 ± 0.033 0.338	0.919 ± 0.003 0.919	0.768 ± 0.063 0.768
Sonnet-4.6	15	0.045 ± 0.010 0.045	0.167 ± 0.011 0.167	0.175 ± 0.018 0.175	0.336 ± 0.042 0.336	0.909 ± 0.004 0.909	0.785 ± 0.054 0.785
GPT-5.4-mini	15	0.083 ± 0.010 0.083	0.182 ± 0.013 0.182	0.246 ± 0.018 0.246	0.353 ± 0.033 0.353	0.918 ± 0.006 0.918	0.766 ± 0.049 0.766
DeepSeek-V4-Pro	15	0.031 ± 0.019 0.031	0.144 ± 0.014 0.144	0.149 ± 0.036 0.149	0.329 ± 0.043 0.329	0.905 ± 0.006 0.905	0.767 ± 0.054 0.767
Qwen3.5-122B	15	0.067 ± 0.011 0.067	0.182 ± 0.017 0.182	0.214 ± 0.027 0.214	0.374 ± 0.027 0.374	0.912 ± 0.008 0.912	0.778 ± 0.058 0.778
DeepSeek-V4-Flash	15	0.015 ± 0.013 0.015	0.123 ± 0.014 0.123	0.114 ± 0.020 0.114	0.257 ± 0.033 0.257	0.896 ± 0.005 0.896	0.698 ± 0.066 0.698
Qwen3.6-35B	15	0.010 ± 0.007 0.010	0.121 ± 0.019 0.121	0.099 ± 0.023 0.099	0.277 ± 0.035 0.277	0.883 ± 0.012 0.883	0.761 ± 0.052 0.761
Qwen3-235B	15	0.014 ± 0.003 0.014	0.143 ± 0.008 0.143	0.114 ± 0.010 0.114	0.311 ± 0.017 0.311	0.892 ± 0.004 0.892	0.776 ± 0.044 0.776
Llama-4	15	0.007 ± 0.007 0.007	0.013 ± 0.013 0.013	0.016 ± 0.016 0.016	0.022 ± 0.022 0.022	0.005 ± 0.005 0.005	0.051 ± 0.051 0.051

dered by inner identifier and grouped by legal area (Civ = civil law, Pub = public law, Crim = criminal law). The heatmap exposes the within-task variance discussed in RQ1: top-tier closed systems hold a tighter per-task spread, while mid-tier and smaller open-weight systems show visibly larger

task-to-task swings for the same overall mean.

E Benchathon participant composition

The Benchathon validation cohort comprises the participants who together produced the human solutions and reviews used throughout this paper.

Table 6: Automatic-metric results on the **ZJS** corpus: per-system bootstrap 95% CIs ($B = 2000$, seed 42) on the surface and embedding-based metrics. Each cell shows the mean stacked above $\pm$ half the width of the 95% CI; dashes mark metrics not computed on this dataset.

System	n	BLEU	ROUGE	METEOR	BERTSc	MovSc	SemSim
Gemini-3.1-Pro	579	0.039 ± 0.003	0.166 ± 0.003	0.152 ± 0.004	0.369 ± 0.007	0.919 ± 0.001	0.780 ± 0.009
GPT-5.4	581	0.078 ± 0.003	0.165 ± 0.002	0.218 ± 0.004	0.350 ± 0.006	0.919 ± 0.001	0.767 ± 0.009
Opus-4.7	581	0.042 ± 0.002	0.159 ± 0.003	0.157 ± 0.004	0.354 ± 0.007	0.913 ± 0.001	0.769 ± 0.009
Sonnet-4.6	581	0.081 ± 0.003	0.164 ± 0.002	0.224 ± 0.005	0.335 ± 0.006	0.917 ± 0.001	0.747 ± 0.009
DeepSeek-V4-Pro	581	0.071 ± 0.003	0.165 ± 0.002	0.208 ± 0.005	0.340 ± 0.006	0.920 ± 0.001	0.762 ± 0.009
GPT-5.4-mini	581	0.031 ± 0.002	0.141 ± 0.002	0.145 ± 0.004	0.324 ± 0.006	0.907 ± 0.001	0.751 ± 0.009
DeepSeek-V4-Flash	581	0.047 ± 0.003	0.155 ± 0.003	0.169 ± 0.004	0.330 ± 0.006	0.909 ± 0.001	0.745 ± 0.009
Gemini-3.1-Flash-Lite	579	0.004 ± 0.001	0.108 ± 0.002	0.073 ± 0.002	0.308 ± 0.005	0.896 ± 0.001	0.762 ± 0.008
Qwen3.5-122B	581	0.018 ± 0.002	0.131 ± 0.002	0.122 ± 0.003	0.317 ± 0.005	0.905 ± 0.001	0.759 ± 0.009
Qwen3.6-35B	581	0.010 ± 0.002	0.116 ± 0.002	0.100 ± 0.002	0.279 ± 0.005	0.899 ± 0.002	0.715 ± 0.009
Llama-4	581	0.004 ± 0.001	0.108 ± 0.003	0.077 ± 0.003	0.297 ± 0.006	0.887 ± 0.001	0.749 ± 0.009
Qwen3-235B	581	0.006 ± 0.001	0.108 ± 0.002	0.080 ± 0.003	0.267 ± 0.005	0.882 ± 0.001	0.724 ± 0.008

Table 7: Automatic-metric and task-correctness results on the **Doctrinal Principles** corpus: per-system bootstrap 95% CIs ($B = 2000$, seed 42) on yes/no (*Ja/Nein*) decision accuracy (Acc.; task-correctness, not a rubric measure) together with the surface and embedding-based metrics. Each cell shows the mean stacked above $\pm$ half the width of the 95% CI; dashes mark metrics not computed on this dataset.

System	n	Acc.	chrF	BLEU	ROUGE	METEOR	BERTSc	MovSc	SemSim
Gemini-3.1-Pro	531	83.4% ± 3.1	0.414 ± 0.007	0.033 ± 0.003	0.164 ± 0.005	0.212 ± 0.008	0.250 ± 0.007	0.840 ± 0.002	0.745 ± 0.008
Opus-4.7	531	81.0% ± 3.3	0.415 ± 0.007	0.030 ± 0.003	0.159 ± 0.005	0.210 ± 0.007	0.247 ± 0.007	0.838 ± 0.002	0.740 ± 0.009
Sonnet-4.6	531	80.6% ± 3.3	0.425 ± 0.006	0.028 ± 0.003	0.147 ± 0.004	0.233 ± 0.006	0.220 ± 0.007	0.843 ± 0.002	0.729 ± 0.009
GPT-5.4	531	78.7% ± 3.4	0.421 ± 0.006	0.031 ± 0.003	0.160 ± 0.005	0.210 ± 0.007	0.249 ± 0.007	0.845 ± 0.002	0.758 ± 0.009
DeepSeek-V4-Pro	531	72.9% ± 3.8	0.419 ± 0.006	0.030 ± 0.003	0.161 ± 0.005	0.210 ± 0.007	0.245 ± 0.008	0.843 ± 0.002	0.752 ± 0.009
Gemini-3.1-Flash-Lite	531	75.5% ± 3.5	0.394 ± 0.007	0.029 ± 0.003	0.162 ± 0.005	0.187 ± 0.008	0.253 ± 0.007	0.838 ± 0.002	0.753 ± 0.009
GPT-5.4-mini	531	70.6% ± 3.8	0.380 ± 0.007	0.028 ± 0.002	0.162 ± 0.005	0.177 ± 0.006	0.250 ± 0.007	0.839 ± 0.002	0.755 ± 0.009
DeepSeek-V4-Flash	530	68.9% ± 3.8	0.412 ± 0.006	0.031 ± 0.003	0.164 ± 0.005	0.209 ± 0.006	0.249 ± 0.007	0.841 ± 0.002	0.749 ± 0.008
Llama-4	531	77.0% ± 3.7	0.388 ± 0.008	0.035 ± 0.003	0.175 ± 0.005	0.199 ± 0.007	0.254 ± 0.008	0.845 ± 0.002	0.747 ± 0.010
Qwen3.5-122B	526	77.6% ± 3.4	0.326 ± 0.009	0.026 ± 0.003	0.165 ± 0.006	0.157 ± 0.008	0.252 ± 0.009	0.827 ± 0.003	0.742 ± 0.009
Qwen3.6-35B	531	70.6% ± 3.6	0.387 ± 0.007	0.027 ± 0.003	0.154 ± 0.005	0.181 ± 0.008	0.239 ± 0.007	0.834 ± 0.002	0.747 ± 0.008
Qwen3-235B	531	72.7% ± 3.8	0.374 ± 0.008	0.025 ± 0.002	0.156 ± 0.005	0.167 ± 0.006	0.241 ± 0.007	0.824 ± 0.002	0.740 ± 0.009

They self-reported their legal-expertise level on a four-stage scale corresponding to standard stages of German legal education and practice: **layperson** (no formal legal training), **law student** (currently enrolled in a German law program, typically working toward the first state examination - *Erstes Staatsexamen*), **graduated** (completed legal studies but not currently in legal practice or *Referen-*

dariat), and **Referendar** (in the approximately two-year practical legal traineeship that follows the first state examination and precedes the second state examination, after which a candidate becomes a fully qualified *Volljurist*). Table 12 breaks the cohort down by self-reported level and working condition (traditional vs. human-AI co-creation), counting both unique participants and annotation contribu-




Figure 3: Per-task per-system mean LLM-judge raw score on the Benchathon subset. Rows are systems ordered by overall mean (top = highest). Columns are the 15 Benchathon tasks ordered by inner_id, grouped by legal area (Civ = civil law, Pub = public law, Crim = criminal law).

Dimension (weight / 100)	Scaffold graded	step
Result correctness / <i>Ergebnis-richtigkeit</i> (20)	<i>Ergebnis</i>	
Doctrinal knowledge / <i>Rechtskenntnis</i> (15)	<i>Definition</i>	
Subsumption quality / <i>Subsumtion</i> (15)	<i>Subsumtion</i>	
Issue identification / <i>Vollständigkeit</i> (10)	<i>Obersatz</i> (issue spotting)	
Legal grounding / <i>Rechtsgrundlagen</i> (10)	<i>Obersatz</i> (norm citation)	
Problem depth / <i>Schwerpunktsetzung</i> (10)	<i>Subsumtion</i> (focal-issue depth)	
Methodological structure / <i>Methodischer Stil</i> (10)	whole-scaffold execution	
Organization / <i>Gliederung</i> (5)	form	
Terminology / <i>Sprache</i> (3)	form	
Formal correctness / <i>Formalia</i> (2)	form	

Table 8: Rubric dimensions mapped to the *Gutachtenstil* scaffold step each grades, with /100 weights from the judge prompt (Appendix L). The judge’s global principle additionally gates all credit on the method: “Keine Punkte für bloße Schlagworte ohne Definition/Subsumtion” (no points for mere buzzwords without definition/subsumption).

Table 9: Per-system mean judge score under the full 10-dimension rubric (*Full*) vs. the seven-dimension reasoning core rescaled to /100 (*Core*), primary judge. System rank agreement full vs. core: Spearman $\rho = 1.00$ (Benchathon), 1.00 (ZJS); Kendall $\tau = 1.00 / 1.00$.

Corpus	System	n	Full	Core
Benchathon	Opus-4.7	15	69.3	67.6
	Gemini-3.1-Pro	15	68.3	66.7
	GPT-5.4	15	66.6	64.6
	Gemini-3.1-Flash-Lite	15	63.5	61.3
	Sonnet-4.6	15	58.9	56.1
	GPT-5.4-mini	15	58.6	55.9
	DeepSeek-V4-Pro	15	56.1	53.2
	Qwen3.5-122B	15	49.4	46.1
	DeepSeek-V4-Flash	15	49.3	45.7
	Qwen3.6-35B	15	42.4	38.4
	Qwen3-235B	15	41.9	38.2
	Llama-4	15	37.8	33.4
ZJS	Gemini-3.1-Pro	578	61.8	59.3
	GPT-5.4	580	60.4	57.8
	Opus-4.7	580	58.7	55.9
	Sonnet-4.6	580	51.5	48.1
	DeepSeek-V4-Pro	580	50.1	46.7
	GPT-5.4-mini	581	49.7	46.2
	DeepSeek-V4-Flash	579	45.9	42.1
	Gemini-3.1-Flash-Lite	578	44.6	40.8
	Qwen3.5-122B	581	44.2	40.3
	Qwen3.6-35B	580	39.8	35.6
	Llama-4	581	36.1	31.7
	Qwen3-235B	581	35.9	31.5

tions per condition.

Table 10: Evaluated LLM systems with provider family, the short alias used throughout the paper, the full model id, tier (flagship / efficiency-oriented / open-weight reference), weights (open vs. closed), reasoning-mode disclosure (explicit thinking variant / no explicit mechanism / undisclosed for closed products with possible internal routing), observed generation temperature, max-output-tokens setting, total generation count on the Benchathon subset, and the count of truncated generations (truncated flag or finish_reason in length / max_tokens) over total. The Google efficiency-tier row aggregates two distinct Google models across corpora: Gemini-3-Flash on Benchathon and Gemini-3.1-Flash-Lite (a separate product line) on ZJS and Doctrinal Principles; see the appendix text.

Provider	Alias	System	Tier	Weights	Reasoning	T	Max out	Gens	Trunc.
Anthropic	Opus-4.7	claude-opus-4-7	Flagship	Closed	undisclosed	1	8000	15	0/15
Anthropic	Sonnet-4.6	claude-sonnet-4-6	Flagship	Closed	undisclosed	1	8000	15	0/15
Google	Gemini-3.1-Flash-Lite	gemini-3.1-flash-lite-preview	Efficiency	Closed	undisclosed	1	8000	15	0/15
Google	Gemini-3.1-Pro	gemini-3.1-pro-preview	Flagship	Closed	undisclosed	1	8000	15	2/15
OpenAI	GPT-5.4	gpt-5.4	Flagship	Closed	undisclosed	1	16000	15	0/15
OpenAI	GPT-5.4-mini	gpt-5.4-mini	Efficiency	Closed	undisclosed	1	8000	15	0/15
Alibaba	Qwen3-235B	Qwen3-235B-A22B-Thinking-2507	Open ref.	Open	explicit (thinking)	0.6	16000	15	0/15
Alibaba	Qwen3.5-122B	Qwen3.5-122B-A10B	Open ref.	Open	none	0.7	8000	15	6/15
Alibaba	Qwen3.6-35B	Qwen3.6-35B-A3B	Open ref.	Open	none	0.7	8000	15	4/15
DeepSeek	DeepSeek-V4-Flash	DeepSeek-V4-Flash	Open ref.	Open	none	1	8000	15	0/15
DeepSeek	DeepSeek-V4-Pro	DeepSeek-V4-Pro	Open ref.	Open	none	1	8000	15	0/15
Meta	Llama-4	Llama-4-Maverick-17B-128E-Instruct-FP8	Open ref.	Open	none	0.6	8000	15	0/15

Table 11: Per-system operational metadata from logged per-call metadata: mean per-generation cost (USD), mean latency (s), and mean output length (words), per corpus. Doctrinal Principles cost was not billed through the logging route and is omitted.

System	B. \$/gen	B. s	B. words	Z. \$/gen	Z. s	Z. words	G. s	G. words
Opus-4.7	\$0.150	73.3	1422	\$0.180	75.2	1710	9.4	104
Sonnet-4.6	\$0.101	104.0	2355	\$0.120	126.1	2796	10.0	167
Gemini-3.1-Flash-Lite	\$0.008	14.9	1281	\$0.003	6.5	699	1.4	79
Gemini-3.1-Pro	\$0.035	102.7	1349	\$0.041	72.6	1558	13.0	95
GPT-5.4	—	70.0	2190	—	79.9	2649	4.6	96
GPT-5.4-mini	—	17.2	1252	—	21.4	1570	2.0	67
Qwen3-235B	\$0.014	125.1	825	\$0.015	148.4	971	46.1	72
Qwen3.5-122B	\$0.018	77.0	1022	\$0.020	90.1	1297	94.2	49
Qwen3.6-35B	\$0.008	74.8	802	\$0.008	64.8	1122	18.5	77
DeepSeek-V4-Flash	\$0.001	89.1	1650	\$0.001	129.1	1906	53.9	101
DeepSeek-V4-Pro	\$0.019	98.6	2012	\$0.023	130.2	2498	16.1	109
Llama-4	\$0.001	37.8	691	\$0.001	37.2	709	9.4	86

Table 12: Benchathon participants by self-reported legal expertise level and working condition (no_ai = traditional, ai = co-creation). The table counts all logged annotation contributions (228); 8 cancelled submissions without participant-written solution text (7 traditional, 1 co-creation) were excluded before evaluation, yielding the 220 solutions (65 traditional, 155 co-creation) analysed throughout the paper.

Expertise	Participants	Trad. annot.	Co-creation annot.
Layperson	5	5	28
Law Student	14	41	60
Graduated	7	10	38
Referendar	7	16	30
Total	33	72	156

F Full Benchathon human-evaluation assignment procedure

The validation subset comprises 45 solutions (15 Benchathon tasks $\times$ 3 provenances: human-traditional, human-AI co-creation, LLM-generated). For each task and condition, the human solution selected is that of the participant whose *expertise proxy* is closest to the median for that legal area and condition. The proxy is a

z-score combining self-assessed competence in the relevant legal area and, where available, the first state examination grade. LLM solutions are drawn through a deterministic rotation across three system tiers (closed flagship, efficiency-oriented, open-weight reference). A per-system load counter spreads selections across systems within each tier, so the validation set covers as many distinct systems as the pool permits. Generations shorter than 200 output tokens are excluded.

Each selected solution receives one review by the task’s creator and three independent blind reviews, each from a distinct grader (180 reviews in total: 45 creator + 135 blind). Task creators do not blind-review solutions in the legal area of their own tasks. In practice, three graders contribute exclusively as blind reviewers across all three legal areas, and the two criminal-law creators additionally contribute blind reviews on civil-law and public-law solutions. The assignment balances reviewer workload and monitors coverage across solution provenance, legal area, and task difficulty. Blind graders are unaware of solution origin: whether a solution was written traditionally, produced in human-AI co-creation, or generated by an LLM. They grade against the same rubric as the LLM judge, at both total and dimension level. Creator reviews are reported as an un-blind reference signal and are not included in inter-rater reliability estimates. The IRR statistics in Table 17 are computed on the full $k=3$ blind-rater pool over all 45 validation solutions.

Roster composition and Calderon ε choice.

The full grading roster is seven people. The $m = 5$ blind subset used in the Calderon blind-pool alt-test are the graders who performed at least one blind review. The remaining two contributed only creator-role reviews and constitute the single-expert reference. Two of the $m = 5$ blind reviewers additionally performed creator reviews on disjoint solution subsets (overlap zero by design), so no solution had the same person in both roles. We use Calderon’s “skilled annotator” cost-benefit penalty $\varepsilon = 0.15$ for the blind-pool procedure, matching the law-students-around-1st-state-exam profile (§5.5). For the single-expert variant we use Calderon’s expert-tier penalty $\varepsilon = 0.20$, since the creator is the most legally expert grader on the roster; at the sample sizes available here, $\varepsilon = 0.15$ and 0.20 yield the same single-expert decision.

G RQ4 supplementary detail

This appendix collects the robustness checks, per-area breakdown, TOST machinery, and blind-human small- n cross-check that support the RQ4 body summary.

Imbalance and skip mechanism. The per-condition imbalance (155 vs. 65 annotations across 33 vs 23 participants) reflects a skip mechanism in the protocol. Participants could de-

cline any item, and less-experienced participants skipped traditional items more often (full per-expertise-stratum breakdown in Appendix E). The traditional pool is therefore expertise-biased upward relative to the co-creation pool. So the headline Δ in the body is a *lower* bound on the per-participant co-creation gain: a forced-completion design would widen the gap, not narrow it.

Robustness across resampling units. The body quotes the participant-clustered bootstrap as the headline. Table 13 reports the same Δ under three additional resampling units, to verify the effect is not an artifact of the chosen unit of inference. All four point estimates agree on sign and within-rounding magnitude. The participant-clustered interval is the widest, correctly recognising that annotations cluster on participant, and is the one we quote in body and Discussion.

Per-area Δ . Per-legal-area Δ values are descriptive only. With three legal areas, $n \leq 53$ per condition per area, and no pre-registered area-level hypothesis, we run no inferential tests of area-level contrasts. We report Table 14 to verify the direction is consistent across all three.

TOST equivalence vs the closed-flagship tier.

We run two one-sided bootstrap tests against a ± 5 -raw-point equivalence band ($\approx \pm 0.9$ grade points). The co-creation mean (65.7) sits -0.1 raw points from the closed-flagship-tier mean (65.8, 95% CI [62.3, 68.9]). One-sided 95% bounds on the difference are $[-3.4, +3.3]$. Both bounds lie inside the ± 5 band. So we reject both $\Delta \leq -5$ and $\Delta \geq +5$: co-creation is statistically equivalent to the closed-flagship tier within the ± 5 -raw-point band.

Blind-human cross-check. On the 15 traditional + 15 co-creation solutions in the RQ5 validation subset ($n=10+11$ participants), the blind-human mean Δ is $+13.6$ raw points (participant-clustered 95% CI $[-1.2, +29.5]$; spans zero at this n). The sign matches the judge result. The gap to the primary GPT-5.4-mini judge’s $+15.7$ -raw-point estimate ($+2.1$ raw points) is small in both absolute and relative terms. It is consistent with GPT-5.4-mini’s near-zero content-dependent calibration offset (Table 21: $\bar{\Delta}_{\text{dir}} \approx 0$ on human-written solutions and ≈ -0.5 on LLM-generated ones, neither significantly different from the pool). The blind-human and LLM-judge estimates of the co-creation gain therefore agree at this sample size, within the inter-rater noise of the blind-pool itself.

Table 13: RQ4 mean difference Δ = co-creation – traditional under four resampling units, raw 0-100 scale. The body quotes the participant-cluster row; the others are reported here for robustness. Bootstrap $B = 2000$ in all cells (seed 7).

Resampling unit	n resample units	Δ	95% CI
i.i.d. annotation	220 annotations	+15.7	[+11.3, +20.0]
Participant cluster	56 participants	+15.7	[+9.6, +22.4]
Task cluster	15 tasks	+15.7	[+11.1, +20.6]
Within-participant paired	23 participants	+13.1	[+8.2, +18.0]

Table 14: RQ4 mean LLM-judge raw scores per legal area, with i.i.d. bootstrap 95% CI on the per-area Δ . Descriptive only; no inferential area-level test is performed.

Legal area	Mean trad	Mean co-cr.	Δ	95% CI	n trad / co-cr.
Civil law	+49.1	+62.9	+13.8	[+5.6, +21.6]	24 / 51
Criminal law	+50.8	+63.9	+13.1	[+8.1, +18.6]	22 / 53
Public law	+50.2	+70.4	+20.2	[+12.2, +28.6]	19 / 51

H Automatic-metric vs. LLM-judge correlation

Table 15 gives the per-generation Pearson r between every automatic metric and the LLM-judge raw score across the three LLM-generation corpora and the two Benchathon human-annotation conditions; Figure 4 repeats the same comparison as a 1×3 panel of per-corpus Pearson and Spearman bar plots (Benchathon | ZJS | Doctrinal Principles).

H.1 Metrics against the human grade

Table 16 re-computes the metric-validity analysis against the mean blind-human grade on the 45-solution validation subset, and adds the residual diagnostic referenced in RQ3: the correlation between (judge – mean blind grade) and each lexical metric. Scored against the human grade, the lexical correlations persist (BLEU $r=+0.70$, ROUGE +0.68, METEOR +0.72 pooled), while sentence-embedding similarity stays weak (+0.32); Mover-Score remains the one embedding exception. The residual correlations are negative for all three lexical metrics (BLEU -0.54, ROUGE -0.44, METEOR -0.43 pooled): the judge rewards high-overlap text *less* than blind human graders do, the opposite of a judge biased toward reference-like style. Grading against the reference solution is, moreover, the shared construct: the blind reviewers grade against the same reference-based rubric (Section 5.5), so some metric-grade correlation is expected for any valid grader. Partial correlations controlling for solution length weaken but do not remove the lexical association. At $n \leq 45$ these numbers indicate direction, not precise magnitude.

I Judge-human agreement statistics

Table 17 reports the headline judge-human and human-human agreement numbers; Figure 5 breaks the same data down by rubric dimension.

The statistic set is non-redundant; each member guards against a failure mode the others cannot detect. Pearson and Spearman measure linear and rank agreement with the mean blind grade; ranks matter because the benchmark’s downstream use is ordering systems. MAE gives the absolute size of judge-human deviations in raw and grade points, which no correlation can show: a judge with a constant strictness offset can hold high r while sitting far from the human scale. Cohen’s κ tests chance-corrected agreement on the one binary decision German grading turns on: pass/fail at *ausreichend*; high continuous correlation does not guarantee agreement at the threshold. ICC(2,1) and ICC(2,k) are not judge statistics at all: they quantify the reliability of a single blind reviewer (the rater the judge would replace) and of the pooled mean (the reference the judge is validated against), the ceiling against which any judge-human agreement must be read. The Calderon alt-test is the primary criterion: it is the one procedure purpose-built to decide whether the judge can stand in for a blind reviewer.

J System prompts for generation

All prompts in this and the following two appendices (Appendices K and L) were sent to the models verbatim in German as reproduced below. Each German prompt is followed by an English translation marked as such; the translations are provided




Figure 4: Per-generation correlation between each automatic metric and the LLM-judge raw score on the Benchathon (left), ZJS (middle), and Doctrinal Principles (right) corpora. Bars report Pearson r and Spearman ρ , with n varying per metric. Metrics are shown in the fixed canonical order shared with the per-corpus correlation table in this appendix. Judge on Doctrinal Principles is the custom 4-dimension rubric.




Figure 5: Per-rubric-dimension agreement between the LLM judge and the mean of the blind human reviewers, split by solution provenance. Left: human-written solutions ($n = 30$). Right: LLM-generated solutions ($n = 15$). Bars: Pearson r (bottom axis). Points: MAE in dimension points (top axis). Dimension ordering is shared across panels (sorted by human-side Pearson). Hatched bars marked n/a indicate that the judge returned a constant score across the 15 LLM-generated solutions on that dimension, leaving the Pearson statistic mathematically undefined (judge variance = 0). The MAE is still well-defined and is plotted.

Table 15: Per-generation Pearson r between each automatic metric and the LLM judge’s raw score across three corpora of LLM generations and the Benchathon human-annotation working conditions (traditional vs. human-AI co-creation). Cell format: r (n). ‘—’ = metric not run on that corpus / condition. Rows are in a fixed canonical order (lexical $\rightarrow$ embedding).

Metric	Benchathon (LLM)	ZJS (LLM)	Grundpr. (LLM)	Bench. human <i>traditional</i>	Bench. human <i>co-creation</i>
BLEU	+0.58 ($n=180$)	+0.42 ($n=6959$)	+0.13 ($n=6365$)	+0.35 ($n=65$)	+0.27 ($n=155$)
ROUGE	+0.69 ($n=180$)	+0.52 ($n=6959$)	+0.25 ($n=6365$)	+0.61 ($n=65$)	+0.37 ($n=155$)
METEOR	+0.61 ($n=180$)	+0.45 ($n=6959$)	+0.28 ($n=6365$)	+0.27 ($n=65$)	+0.36 ($n=155$)
chrF	—	—	+0.33 ($n=6365$)	—	—
BERTScore	+0.48 ($n=180$)	+0.37 ($n=6959$)	+0.26 ($n=6365$)	—	—
MoverScore	+0.62 ($n=180$)	+0.49 ($n=6955$)	+0.12 ($n=6365$)	+0.38 ($n=16$)	+0.61 ($n=28$)
Semantic sim.	+0.28 ($n=180$)	+0.19 ($n=6959$)	+0.15 ($n=6365$)	+0.20 ($n=16$)	+0.37 ($n=29$)

Table 16: Automatic metrics correlated with the mean blind-human raw grade on the 45-solution validation subset (Pearson r ; 30 human-written, 15 LLM-generated solutions). The residual column correlates (judge – mean blind grade) with the metric on the pooled subset; the uniformly negative values run opposite to a judge reference-style bias. chrF is not computed on Benchathon; BERTScore is unavailable for the human-written solutions. Small- n cells indicate direction, not precise magnitude. Bootstrap CIs and per-split residuals in `metric_vs_human.json` of the released pipeline.

Metric	Human-writ.	LLM-gen.	Pooled	Residual (pooled)
BLEU	+0.70 ($n=30$)	+0.70 ($n=15$)	+0.70 ($n=45$)	-0.54 ($n=45$)
ROUGE	+0.67 ($n=30$)	+0.68 ($n=15$)	+0.68 ($n=45$)	-0.44 ($n=45$)
METEOR	+0.71 ($n=30$)	+0.78 ($n=15$)	+0.72 ($n=45$)	-0.43 ($n=45$)
BERTScore	—	+0.59 ($n=15$)	+0.59 ($n=15$)	—
MoverScore	+0.87 ($n=7$)	+0.85 ($n=15$)	+0.85 ($n=22$)	—
Semantic sim.	+0.53 ($n=7$)	-0.01 ($n=15$)	+0.32 ($n=22$)	—

for reader convenience and were not used in any model interaction.

J.1 Exam-style cases (Klausuren)

Du bist ein hochqualifizierter juristischer Bearbeiter für deutsches Recht. Du bearbeitest juristische Klausuren methodisch sauber, präzise und strukturiert nach den Standards deutscher juristischer Ausbildung und Prüfungspraxis.

Du argumentierst streng fallbezogen und orientierst dich an den Anforderungen universitärer Klausuren sowie des ersten Staatsexamens. Du löst Aufgaben eigenständig und ohne Meta-Kommentare.

English translation (reader reference; not sent to any model):

You are a highly qualified legal practitioner for German law. You work through legal exam-style cases (*Klausuren*) methodically, precisely, and in a structured manner, in line with the standards of German legal education and examination practice.

You argue strictly with reference to the case at hand and orient yourself to the requirements of university exam-style cases and of the first state examination (*Erstes Staatsexamen*). You solve tasks independently and without meta-commentary.

J.2 Doctrinal Principles (Grundprinzipien)

Du bist ein hochqualifizierter juristischer Bearbeiter für deutsches Recht. Du beantwortest juristische Fragen präzise, methodisch sauber und streng fallbezogen nach den Standards deutscher Rechtswissenschaft.

Deine Antworten sollen knapp, klar und juristisch korrekt sein. Im Mittelpunkt steht die zutreffende Anwendung rechtlicher Grundprinzipien auf den konkreten Sachverhalt. Du argumentierst eigenständig und ohne Meta-Kommentare.

English translation (reader reference; not sent to any model):

You are a highly qualified legal practitioner for German law. You answer legal questions precisely, methodically, and strictly with reference to the case at hand, in line with the standards of German legal scholarship.

Your answers should be concise, clear, and legally correct. The focus is on the accurate application of fundamental legal principles to the concrete fact pattern. You argue independently and without meta-commentary.

K Instruction prompts for generation

K.1 Exam-style cases (Klausuren)

Bearbeite die folgende juristische Aufgabe wie eine echte deutsche Jura-Klausur.

Ziel der Bearbeitung Erstelle eine vollständige juristische Falllösung nach den methodischen Standards deutscher Rechtswissenschaft. Im Mittelpunkt steht subsumptionsbasierte Rechtsanwendung.

Die Bearbeitung soll zeigen, dass du:

- die rechtlichen Kernprobleme des Falls erkennt,
- die relevanten Anspruchsgrundlagen bzw. Prüfungsprogramme auswählt,
- die Prüfung methodisch korrekt strukturiert,
- Definitionen und Streitstände präzise einordnet,
- vor allem aber die konkreten Tatsachen des Falls sauber unter die rechtlichen Voraussetzungen subsumierst.

Die Qualität der juristischen Argumentation ist wichtiger als reine Vollständigkeit oder Länge.

Stil, Methodik und Gliederung Arbeite im Stil einer deutschen juristischen Klausur.

Grundsätzlich ist die Lösung im Gutachtenstil zu verfassen. Das bedeutet insbesondere:

- Formulierung eines Obersatzes,
- Definition der rechtlichen Voraussetzungen,
- Subsumtion anhand des konkreten Sachverhalts,
- nachvollziehbares Zwischenergebnis.

Nur wenn aus Aufgabenstellung, Ausbildungsniveau oder Bearbeitungskontext erkennbar ist, dass eine Referendarklausur oder ein Urteil verlangt wird, ist stattdessen im Urteilsstil zu arbeiten. In diesem Fall soll die Lösung den typischen Aufbau gerichtlicher Entscheidungen widerspiegeln, insbesondere eine ergebnisorientierte Darstellung mit anschließender Begründung.

Die juristische Methodik muss unabhängig vom Stil jederzeit klar erkennbar bleiben.

Die Lösung ist klausurtypisch und hierarchisch sauber zu gliedern. Nutze die in der deutschen Juristenausbildung üblichen Gliederungsebenen und Bezeichnungen.

Typischerweise:

- A., B., C. für Hauptabschnitte,
- I., II., III. für Unterabschnitte,
- 1., 2., 3. für weitere Ebenen,
- a), b), c) für Unterpunkte,
- aa), bb), cc) für vertiefte Untergliederungen.

Jede Gliederungsebene soll inhaltlich sinnvoll, logisch konsistent und leicht nachvollziehbar sein. Das bedeutet insbesondere:

- strukturiere die Lösung nach Anspruchsgrundlagen, Rechtsbehelfen oder Prüfungsprogrammen,
- prüfe Anspruchsvoraussetzungen in methodisch sinnvoller Reihenfolge,
- bilde erkennbare Prüfungsschwerpunkte heraus,
- nutze Zwischenergebnisse zur Strukturierung der weiteren Prüfung,
- vermeide unstrukturierte Textblöcke oder springhafte Argumentation.

Die Qualität der juristischen Strukturierung ist Teil der Bearbeitungsleistung.

Schwerpunktsetzung Gewichte die Bearbeitung klausurtypisch.

Das bedeutet:

- konzentriere dich auf die rechtlich schwierigen und problematischen Punkte,
- behandle offensichtliche oder unproblematische Voraussetzungen kurz,
- erkenne versteckte Probleme und Schwerpunktsetzungen im Sachverhalt,
- vermeide lange Ausführungen zu irrelevanten Nebenaspekten.

Umgang mit Streitständen Die Darstellung relevanter Streitstände ist in juristischen Klausuren grundsätzlich erwünscht, insbesondere wenn sie zu den Kernproblemen des Falls gehören oder methodisch vertiefte Argumentation ermöglichen.

Wenn mehrere vertretbare Ansichten bestehen:

- stelle die wesentlichen Positionen präzise dar,
- arbeite Unterschiede in Argumentation oder dogmatischer Herleitung heraus,
- argumentiere methodisch und fallbezogen,
- entscheide dich nachvollziehbar für eine Ansicht,
- arbeite anschließend konsequent mit der gewählten Ansicht weiter.

Die Darstellung darf über das absolut unmittelbar Notwendige hinausgehen, sofern sie klausurtypisch ist und zum Schwerpunkt der Aufgabe passt. Vermeide jedoch:

- rein auswendig gelernte oder schematische Streiddarstellungen ohne Fallbezug,
- extensive Meinungsstreite zu nebensächlichen Problemen,
- reine Namens- oder Schlagwortaufzählungen ohne argumentative Einordnung.

Umgang mit dem Sachverhalt Arbeite streng sachverhaltsbezogen.

- Nutze konkrete Tatsachen aktiv in der Subsumtion.
- Erfinde keine zusätzlichen Tatsachen.
- Unterstelle keine Informationen, die nicht im Sachverhalt enthalten sind.
- Ignoriere keine auffälligen Hinweise oder Indizien.

Wenn Informationen fehlen, arbeite mit dem vorhandenen Sachverhalt weiter und kennzeichne verbleibende Unsicherheiten juristisch sauber.

Normen und Zitierweise Nenne die einschlägigen gesetzlichen Grundlagen präzise.

- Zitiere Normen möglichst vollständig.
- Nutze Absätze, Sätze, Nummern oder Buchstaben, wenn relevant.
- Verwende deutsche juristische Terminologie.

Beispiele:

- „§ 823 Abs. 1 BGB“
- „Art. 12 Abs. 1 GG“
- „§ 242 StGB“

Sprachstil

- Schreibe sachlich, präzise und juristisch.
- Vermeide Meta-Kommentare.
- Erwähne nicht, dass du ein KI-System bist.
- Gib keine allgemeinen Haftungsausschlüsse.
- Verweise nicht auf externe Beratung.
- Antworte direkt mit der juristischen Lösung.

Umfang Die Lösung soll so ausführlich sein, wie es für eine gute bis sehr gute juristische Klausurbearbeitung erforderlich ist.

- Wichtige Probleme sollen vertieft behandelt werden.
- Unproblematische Punkte dürfen knapp bleiben.
- Die Lösung soll nicht künstlich verlängert werden.

Ausgabeformat

- Gib ausschließlich die juristische Lösung aus.
- Nutze sinnvolle Überschriften und Gliederungsebenen.
- Verwende Fließtext.
- Keine Stichpunkte außer wenn methodisch sinnvoll.
- Keine zusätzlichen Erklärungen außerhalb der Falllösung.

Aufgabe {{sachverhalt}}

English translation (reader reference; not sent to any model):

Work through the following legal task as if it were a genuine German law exam-style case (*Klausur*).

Goal of the work Produce a complete legal case solution (*Falllösung*) according to the methodological standards of German legal scholarship. The focus is on subsumption-based (*subsumtionsbasiert*) application of the law.

The work should demonstrate that you:

- identify the core legal problems of the case,
- select the relevant causes of action or examination programmes,
- structure the analysis methodologically correctly,
- classify definitions and scholarly debates (*Streitstände*) precisely,
- and, above all, cleanly subsume the concrete facts of the case under the legal requirements.

The quality of the legal argumentation is more important than mere completeness or length.

Style, methodology, and structure Work in the style of a German legal exam-style case.

As a general rule, the solution is to be written in the assessment style (*Gutachtenstil*). This means in particular:

- formulation of a major premise (*Obersatz*),
- definition of the legal requirements,
- subsumption on the basis of the concrete fact pattern,
- a comprehensible intermediate result.

Only when it is evident from the task, the level of training, or the work context that a trainee's exam-style case (*Referendarsklausur*) or a judgement is required should the judgement style (*Urteilstil*) be used instead. In that case, the solution should reflect the typical structure of judicial decisions, in particular a result-oriented presentation followed by the reasoning.

Regardless of style, the legal methodology must remain clearly recognizable at all times.

The solution must be structured in a typical exam-style way, with a clean hierarchy. Use the levels and labels customary in German legal education.

Typically:

- A., B., C. for main sections,
- I., II., III. for subsections,
- 1., 2., 3. for further levels,
- a), b), c) for sub-points,
- aa), bb), cc) for deeper sub-structures.

Each level of structure should be substantively meaningful, logically consistent, and easy to follow. This means in particular:

- structure the solution by causes of action, legal remedies, or examination programmes,

- examine the requirements of a claim in a methodologically sensible order,
- bring out recognizable focal points of the analysis,
- use intermediate results to structure the further analysis,
- avoid unstructured blocks of text or erratic argumentation.

The quality of the legal structuring is part of the work being assessed.

Setting priorities Weight the work in the manner typical of an exam-style case.

This means:

- concentrate on the legally difficult and problematic points,
- treat obvious or unproblematic requirements briefly,
- recognise hidden problems and focal points in the fact pattern,
- avoid lengthy discussion of irrelevant secondary aspects.

Handling scholarly debates The presentation of relevant scholarly debates (*Streitstände*) is generally desirable in legal exam-style cases, especially when they concern the core problems of the case or enable methodologically deeper argumentation.

When several defensible views exist:

- present the essential positions precisely,
- bring out differences in argumentation or doctrinal derivation,
- argue methodologically and with reference to the case,
- decide for one view in a comprehensible way,
- and then work consistently with the chosen view.

The presentation may go beyond what is absolutely strictly necessary, provided it is typical of an exam-style case and fits the focus of the task. However, avoid:

- purely rote or schematic presentations of legal controversies without a connection to the case,
- extensive doctrinal disputes on secondary issues,
- mere name-dropping or buzzword listings without argumentative classification.

Handling the fact pattern Work strictly with reference to the fact pattern (*Sachverhalt*).

- Use concrete facts actively in the subsumption.
- Do not invent additional facts.
- Do not assume information that is not contained in the fact pattern.
- Do not ignore any conspicuous indications or pieces of evidence.

If information is missing, continue working with the fact pattern as given and mark any remaining uncertainties in a legally clean manner.

Norms and citation style Cite the applicable statutory grounds precisely.

- Cite norms as completely as possible.
- Use paragraphs, sentences, numbers, or letters when relevant.
- Use German legal terminology.

Examples:

- “§ 823 Abs. 1 BGB”
- “Art. 12 Abs. 1 GG”
- “§ 242 StGB”

Language and style

- Write objectively, precisely, and in a legal register.
- Avoid meta-commentary.
- Do not mention that you are an AI system.
- Do not give general disclaimers.
- Do not refer to external advice.
- Respond directly with the legal solution.

Length The solution should be as detailed as required for a good to very good legal exam-style case solution.

- Important problems should be treated in depth.
- Unproblematic points may be kept short.
- The solution should not be artificially extended.

Output format

- Output only the legal solution.
- Use meaningful headings and levels of structure.
- Use running prose.
- No bullet points except where methodologically sensible.
- No additional explanations outside the case solution.

Task {{sachverhalt}}

K.2 Doctrinal Principles (Grundprinzipien)

Bearbeite die folgende juristische Frage nach deutschem Recht.

Ziel Gib eine kurze, präzise und methodisch saubere juristische Antwort auf die konkrete Fragestellung.

Die Antwort soll:

- die rechtlich relevante Kernfrage erkennen,
- die einschlägigen Normen oder Prinzipien benennen,
- den konkreten Sachverhalt knapp subsumieren,
- und zu einer klaren Ja-/Nein-Entscheidung gelangen.

Stil Die Antwort soll deutlich kürzer und fokussierter sein als eine vollständige Klausurlösung.

- Keine umfangreiche Gliederung.
- Keine langen Streitstandsdebatten.
- Keine künstliche Verlängerung.
- Keine allgemeinen Lehrbuchausführungen.

Wenn mehrere vertretbare Ansichten bestehen:

- stelle die wesentlichen Unterschiede knapp dar,
- entscheide dich nachvollziehbar,
- und arbeite konsistent mit der gewählten Ansicht weiter.

Sachverhaltsbezug Arbeite streng sachverhaltsbezogen.

- Nutze nur Informationen aus Fall und Fragestellung.
- Erfinde keine zusätzlichen Tatsachen.
- Vermeide abstrakte Standardformulierungen ohne Fallbezug.

Ausgabeformat Gib ausschließlich gültiges JSON im folgenden Format aus:

```
{
  "answer": "Ja" oder "Nein",
  "reasoning": "kurze juristische Begründung"
}
```

Fall {{fall}}

Fragestellung {{task}}

English translation (reader reference; not sent to any model):

Work through the following legal question under German law.

Goal Give a short, precise, and methodologically clean legal answer to the concrete question.

The answer should:

- identify the legally relevant core question,
- name the applicable norms or principles,
- briefly subsume the concrete fact pattern,
- and arrive at a clear Ja/Nein decision.

Style The answer should be considerably shorter and more focused than a full exam-style case solution.

- No extensive structuring.
- No long discussions of scholarly debates.
- No artificial padding.
- No general textbook expositions.

When several defensible views exist:

- briefly present the essential differences,
- decide for one view in a comprehensible way,
- and work consistently with the chosen view.

Reference to the fact pattern Work strictly with reference to the fact pattern.

- Use only information from the case and the question.
- Do not invent additional facts.
- Avoid abstract standard formulations without a connection to the case.

Output format Output only valid JSON in the following format:

```
{
  "answer": "Ja" or "Nein",
  "reasoning": "short legal justification"
}
```

Case {{fall}}

Question {{task}}

L Evaluation prompts for the LLM judge

L.1 Exam-style cases (Klausuren)

Du bist ein strenger, aber fairer Korrektor ("LLM Judge") für juristische Klausuren deutscher Studierender bis zum 1. Staatsexamen.

Du bewertest ausschließlich auf Basis von: 1. Angabe/Sachverhalt, 2. Studierendenlösung, 3. Musterlösung.

Nutze keine externen Quellen und erfinde keine Tatsachen.

Grundsätze:

- Musterlösung ist Referenz. Vergib aber vergleichbar Punkte für vertretbare Alternativlösungen, wenn sie mit der Angabe vereinbar und methodisch sauber begründet sind.
- Folgefehlerprinzip: Einen Fehler nicht doppelt bestrafen. Wenn ein früher Fehler spätere Teile beeinflusst, bewerte die spätere Argumentation unter der

Prämisse des studentischen Zwischenergebnisses, soweit methodisch sauber.

- Keine Punkte für bloße Schlagworte ohne Definition/Subsumtion.
- Beziehe dich in der Begründung konkret darauf, was in der Studierendenlösung steht und wie es zur Musterlösung passt.

Bewertungsschema (insgesamt 100 Punkte):

1. Ergebnisrichtigkeit (0-20): **Vollpunkte (≈18-20):** Kernergebnisse stimmen mit Musterlösung überein oder sind vertretbar gleichwertig (z.B. anderes, aber methodisch korrekt begründetes Ergebnis). **Mittel (≈10-14):** Haupttendenz stimmt, aber ein wesentliches Ergebnis ist falsch oder Rechtsfolgen sind merklich verfehlt. **Niedrig (≈0-8):** Zentrale Ergebnisse verfehlt/inkonsistent; Lösung läuft am Fall vorbei.
2. Vollständigkeit & Problemidentifikation (0-10): **Vollpunkte (≈9-10):** Alle wesentlichen Probleme aus Angabe/Musterlösung erkannt; keine großen „blinden Flecken“. **Mittel (≈5-7):** Ein wesentliches Problem fehlt oder ein Schwerpunkt wird übersehen; sonst ordentlich. **Niedrig (≈0-4):** Mehrere zentrale Probleme fehlen oder es werden viele irrelevante Nebenthemen geprüft („Themenklausur“ statt Falllösung).
3. Rechtsgrundlagen & Prüfungssystematik (0-10): **Vollpunkte (≈9-10):** Richtige Anspruchsgrundlagen/Prüfungsprogramme, logische Reihenfolge, saubere Prüfungsabschnitte. **Mittel (≈5-7):** Kleine Systematikfehler (Reihenfolge, einzelne falsche/unnötige Grundlagen), aber Gesamtstruktur trägt. **Niedrig (≈0-4):** Falsche Grundstruktur (z.B. falscher Rechtsbehelf/Anspruch, falscher Prüfungsrahmen), sodass die Prüfung methodisch entgleist.
4. Rechtskenntnis (Definitionen, Normen, Streitstände) (0-15): **Vollpunkte (≈13-15):** Definitionen korrekt, Normen passend, Streitstände nur dort, wo nötig; Streitentscheid begründet. **Mittel (≈8-11):** Einzelne Definitionen/Normen ungenau; Streitstände schematisch oder lückenhaft, aber nicht fallentscheidend falsch. **Niedrig (≈0-7):** Häufig falsche Definitionen/Normen oder erfundene Anforderungen; Streitstände wirr/fehlerhaft.
5. Subsumtion & Argumentationsqualität (Fallbezug) (0-15): **Vollpunkte (≈13-15):** Konsequenter Tatsachenbezug, saubere Subsumtion, nachvollziehbare Argumente, Abwägungen strukturiert. **Mittel (≈8-11):** Teilweise nur abstrakt/„leerformelhaft“, aber wesentliche Punkte werden noch fallbezogen gelöst. **Niedrig (≈0-7):** Kaum Subsumtion, überwiegend Definitionen ohne Anwendung; Ergebnisse nicht begründet oder widersprüchlich.
6. Schwerpunktsetzung & Problemtiefe (0-10): **Vollpunkte (≈9-10):** Langer Sachverhaltsblock/Indizien → angemessen vertieft; einfache Punkte kurz; Schwerpunkt stimmt mit Musterlösung und „Signalen“ der Angabe überein. **Mittel (≈5-7):** Entweder Überlänge bei Nebensachen oder Untergewichtung eines Schwerpunkts, aber noch erkennbares Klausurbewusstsein. **Niedrig (≈0-4):** Massive Fehlgewichtung: Hauptproblem kaum behandelt, Nebensachen dominieren.
7. Methodischer Stil: Obersatz → Definition → Subsumtion → Ergebnis (0-10): **Vollpunkte (≈9-10):** Bei allen wichtigen Prüfungspunkten klar erkennbar; Zwischenergebnisse werden gezogen;

auch im Urteilsstil sinngemäß eingehalten. **Mittel (≈5-7):** Grundschemata vorhanden, aber häufiger vermischt/ausgelassen; Definitionen oder Ergebnisse fehlen gelegentlich. **Niedrig (≈0-4):** Kein prüfungsorientierter Stil; nur Nacherzählung, Stichworte, oder reine Ergebnisbehauptungen.

8. Gliederung, Gliederungsebenen & Leseführung (0-5): **5:** Stringente Gliederung, Ebenen sauber, Prüfpunkte auffindbar. **3-4:** Kleinere Inkonsistenzen (Ebenensprünge, Überschriften nicht passgenau). **0-2:** Unstrukturierter Textblock; Gliederung irreführend.
9. Sprache & juristische Terminologie (0-3): **3:** Präzise, juristisch sauber, gut verständlich. **2:** Teilweise unständig/ungenau, aber verständlich. **0-1:** Häufig missverständlich, falsche Begriffe, viele sprachliche Brüche.
10. Formalia: Normzitierweise, Sachverhaltsarbeit, Sorgfalt (0-2): **2:** Normen überwiegend korrekt zitiert; Sachverhalt korrekt verarbeitet; keine groben Schnitzer. **1:** Einzelne Zitier-/Sorgfaltsfehler, nicht gravierend. **0:** Häufige grobe Fehler (Rollen vertauscht, zentrale Fakten falsch wiedergegeben, Normchaos).

Umrechnung Rohpunkte → Notenpunkte (0-18; <4 = nicht bestanden): 0-12→0, 13-25→1, 26-38→2, 39-49→3, 50-53→4, 54-56→5, 57-59→6, 60-63→7, 64-66→8, 67-69→9, 70-73→10, 74-76→11, 77-79→12, 80-83→13, 84-86→14, 87-89→15, 90-93→16, 94-96→17, 97-100→18

Ausgabeformat:

- Gib eine JSON-Ausgabe ohne zusätzlichen Text aus.
- Für jede Dimension: score (0 bis max, ggf. 0.5 Schritte), max, und eine Begründung als zusammenhängender Absatz (3-7 Sätze).
- Danach: total_score (0-100), grade_points (0-18), passed (true/false), kurze Gesamtwürdigung (2-5 Sätze) und 3 konkrete Verbesserungshinweise.

ANGABE / SACHVERHALT:
{{{ANGABE_TEXT}}}
MUSTERLÖSUNG: {{{MUSTERLÖSUNG_TEXT}}}
STUDIENDENLÖSUNG: {{{STUDIENDENLÖSUNG_TEXT}}}

Aufgabe: Bewerte die Lösung nach dem vorgegebenen 100-Punkte-Schema. Vergib zu jeder Dimension Punkte und schreibe eine Begründung als Absatz. Berechne die Summe (0-100) und wandle in Notenpunkte (0-18) um. Markiere passed=true nur, wenn grade_points >= 4.

Gib ausschließlich JSON aus.

English translation (reader reference; not sent to any model):

You are a strict but fair grader (“LLM Judge”) for legal exam-style cases written by German students up to the first state examination (*Erstes Staatsexamen*).

You evaluate exclusively on the basis of: 1. the task/fact pattern, 2. the student solution (*Studierendenlösung*), 3. the reference solution (*Musterlösung*).

Do not use any external sources and do not invent any facts.

Principles:

- The reference solution is the benchmark. However, award comparable points for defensible alternative solutions, provided they are compatible with the task and methodologically soundly reasoned.

- Follow-on-error principle (*Folgefehlerprinzip*): do not punish one error twice. If an early error affects later parts, evaluate the later argumentation on the assumption of the student's intermediate result, to the extent this is methodologically clean.
- No points for mere buzzwords without definition / subsumption.
- In your justification, refer concretely to what the student solution says and how it fits the reference solution.

Grading scheme (100 points in total):

1. Correctness of result (0-20): **Full points (≈18-20)**: core results match the reference solution or are defensibly equivalent (e.g. a different but methodologically correctly reasoned result). **Middle (≈10-14)**: the main tendency is right, but one essential result is wrong or the legal consequences are noticeably off. **Low (≈0-8)**: central results are missed / inconsistent; the solution misses the point of the case.
2. Completeness and problem identification (0-10): **Full points (≈9-10)**: all essential problems from the task / reference solution are recognised; no major "blind spots". **Middle (≈5-7)**: one essential problem is missing or a focal point is overlooked; otherwise sound. **Low (≈0-4)**: several central problems are missing, or many irrelevant side topics are examined ("topic essay" instead of case solution).
3. Legal bases and examination systematics (0-10): **Full points (≈9-10)**: correct causes of action / examination programmes, logical sequence, clean examination sections. **Middle (≈5-7)**: small systematic errors (sequence, individual incorrect / unnecessary bases), but the overall structure holds. **Low (≈0-4)**: wrong basic structure (e.g. wrong legal remedy / claim, wrong examination framework), so that the analysis methodologically derails.
4. Legal knowledge (definitions, norms, scholarly debates) (0-15): **Full points (≈13-15)**: definitions are correct, norms are apt, scholarly debates appear only where needed; the decision in the debate is reasoned. **Middle (≈8-11)**: individual definitions / norms are imprecise; scholarly debates are schematic or patchy, but not wrong in a case-decisive way. **Low (≈0-7)**: frequently wrong definitions / norms or invented requirements; scholarly debates are muddled / faulty.
5. Subsumption and quality of argumentation (connection to the case) (0-15): **Full points (≈13-15)**: consistent reference to facts, clean subsumption, comprehensible arguments, structured balancing. **Middle (≈8-11)**: in part only abstract / "empty-formula-like", but essential points are still solved with reference to the case. **Low (≈0-7)**: hardly any subsumption, predominantly definitions without application; results are unreasoned or contradictory.
6. Setting priorities and depth of analysis (0-10): **Full points (≈9-10)**: long fact-pattern blocks / indications are appropriately treated in depth; easy points are kept short; the focus matches the reference solution and the "signals" of the task. **Middle (≈5-7)**: either overlength on minor matters or under-weighting of a focal point, but still recognizable exam awareness. **Low (≈0-4)**: massive misallocation: the main problem is barely treated, secondary matters dominate.
7. Methodological style: major premise → definition → subsumption → result (0-10): **Full points (≈9-10)**: clearly recognizable at all important examination points; intermediate results are drawn; also observed in spirit when written in judgement style.

Middle (≈5-7): the basic scheme is present but is frequently blended or omitted; definitions or results are occasionally missing. **Low (≈0-4)**: no examination-oriented style; only narration, bullet points, or bare assertions of results.

8. Structure, structural levels, and reader guidance (0-5): **5**: stringent structure, clean levels, examination points are findable. **3-4**: small inconsistencies (level jumps, headings not precisely fitting). **0-2**: unstructured block of text; structure is misleading.
9. Language and legal terminology (0-3): **3**: precise, legally clean, easily understandable. **2**: partly clumsy / imprecise, but understandable. **0-1**: frequently misleading, wrong terms, many linguistic breaks.
10. Formalia: norm citation style, work with the fact pattern, care (0-2): **2**: norms cited predominantly correctly; fact pattern processed correctly; no gross blunders. **1**: isolated citation / care errors, not serious. **0**: frequent gross errors (roles swapped, central facts misreported, norm chaos).

Conversion of raw points → grade points (0-18; <4 = fail): 0-12→0, 13-25→1, 26-38→2, 39-49→3, 50-53→4, 54-56→5, 57-59→6, 60-63→7, 64-66→8, 67-69→9, 70-73→10, 74-76→11, 77-79→12, 80-83→13, 84-86→14, 87-89→15, 90-93→16, 94-96→17, 97-100→18

Output format:

- Output a JSON response without any additional text.
- For each dimension: score (0 to max, in 0.5 steps if needed), max, and a justification as a coherent paragraph (3-7 sentences).
- Then: `total_score` (0-100), `grade_points` (0-18), `passed` (true/false), a short overall assessment (2-5 sentences), and 3 concrete suggestions for improvement.

TASK / FACT PATTERN: {{{ANGABE_TEXT}}}
REFERENCE SOLUTION: {{{MUSTERLOESUNG_TEXT}}}
STUDENT SOLUTION: {{{LOESUNG_TEXT}}}

Task: Evaluate the solution according to the prescribed 100-point scheme. Award points for each dimension and write a justification as a paragraph. Calculate the sum (0-100) and convert it into grade points (0-18). Set `passed=true` only if `grade_points >= 4`.

Output only JSON.

L.2 Doctrinal Principles (Grundprinzipien)

Du bist ein strenger, aber fairer juristischer Korrektor für kurze juristische Begründungsfragen nach deutschem Recht.

Du bewertest ausschließlich auf Basis von:

1. Fall
2. Fragestellung
3. Referenzlösung
4. Zu bewertende Antwort

Nutze keine externen Quellen und erfinde keine Tatsachen.

Bewertungsgrundsätze

- Die Referenzlösung ist Orientierung, nicht die einzig mögliche vertretbare Lösung.
- Vergib volle oder nahezu volle Punkte auch für methodisch vertretbare Alternativbegründungen.

- Entscheidend ist die juristische Richtigkeit und Qualität der Begründung, nicht die sprachliche Ähnlichkeit zur Referenz.
- Keine Punkte für bloße Schlagworte ohne nachvollziehbare Begründung.
- Eine falsche Ja/Nein-Entscheidung begrenzt die Maximalpunktzahl erheblich, selbst wenn Teile der Begründung juristisch plausibel sind.

Bewertungsschema (100 Punkte)

1. Ergebnisrichtigkeit (0-40)

- Ist die Ja/Nein-Entscheidung juristisch zutreffend oder vertretbar?

1. Rechtskenntnis und Normbezug (0-25)

- Werden einschlägige Normen, Definitionen oder Prinzipien korrekt erkannt und verwendet?

1. Subsumtion und Fallbezug (0-25)

- Wird der konkrete Sachverhalt nachvollziehbar unter die rechtlichen Voraussetzungen subsumiert?

1. Klarheit und Präzision (0-10)

- Ist die Antwort präzise, verständlich und methodisch sauber formuliert?

Ausgabeformat Gib ausschließlich gültiges JSON aus:

```
{
  "scores": {
    "result_correctness": {
      "score": ...,
      "max": 40,
      "reason": "..."
    },
    "legal_knowledge": {
      "score": ...,
      "max": 25,
      "reason": "..."
    },
    "subsumption": {
      "score": ...,
      "max": 25,
      "reason": "..."
    },
    "clarity": {
      "score": ...,
      "max": 10,
      "reason": "..."
    }
  },
  "total_score": ...,
  "overall_assessment": "..."
}
```

Fall {{fall}}

Fragestellung {{task}}

Referenzlösung Antwort: {{binary_solution}}

Begründung: {{reasoning}}

Zu bewertende Antwort {{answer}}

English translation (reader reference; not sent to any model):

You are a strict but fair legal grader for short legal reasoning questions under German law.

You evaluate exclusively on the basis of:

1. the case,
2. the question,
3. the reference solution,
4. the answer to be evaluated.

Do not use any external sources and do not invent any facts.

Grading principles

- The reference solution is a guide, not the only defensible solution.
- Award full or near-full points also for methodologically defensible alternative justifications.
- What matters is the legal correctness and quality of the reasoning, not linguistic similarity to the reference.
- No points for mere buzzwords without a comprehensible justification.
- A wrong Ja/Nein decision substantially limits the maximum score, even if parts of the reasoning are legally plausible.

Grading scheme (100 points)

1. Correctness of result (0-40)

- Is the Ja/Nein decision legally correct or defensible?

1. Legal knowledge and reference to norms (0-25)

- Are the applicable norms, definitions, or principles correctly identified and used?

1. Subsumption and reference to the case (0-25)

- Is the concrete fact pattern subsumed under the legal requirements in a comprehensible way?

1. Clarity and precision (0-10)

- Is the answer precisely, understandably, and methodologically cleanly formulated?

Output format Output only valid JSON:

```
{
  "scores": {
    "result_correctness": {
      "score": ...,
      "max": 40,
      "reason": "..."
    },
    "legal_knowledge": {
      "score": ...,
      "max": 25,
      "reason": "..."
    },
    "subsumption": {
      "score": ...,
      "max": 25,
      "reason": "..."
    },
    "clarity": {
      "score": ...,
      "max": 10,
      "reason": "..."
    }
  },
  "total_score": ...,
  "overall_assessment": "..."
}
```

Case {{fall}}

Question {{task}}

Reference solution Answer: {{binary_solution}}
Reasoning: {{reasoning}}

Answer to be evaluated {{answer}}

M LLM-judge configuration

The primary judge GPT-5.4-mini is the one used for every reported leaderboard score and for the RQ5 within-judge stochasticity analysis ($k=3$ repeats per generation). The five secondary judges Opus-4.7, Gemini-3.1-Pro, DeepSeek-V4-Pro, Qwen3.5-397B-A17B, and Sonnet-4.6 evaluate every Benchathon cell once for the RQ5 between-judge bias analysis. The full model ids are shown in the *Model* column below; see Table 10 (Appendix B) for the canonical alias-to-id mapping.

The full judge prompts are reproduced in Appendix L: the exam-case rubric and the Doctrinal Principles rubric. All six judges receive the same prompt and the same input ordering (task, reference solution, solution to be evaluated); only the underlying model changes.

N Judge calibration - supplementary detail

This appendix collects the supplementary numerical detail behind the RQ5 per-judge calibration analysis: the judge-swapped system leaderboard and its rank agreement in Tables 19 and 20, the six-judge calibration table in Table 21, the within-judge per-pass breakdown of the primary GPT-5.4-mini judge in Table 22, the per-annotator alt-test p -values, a robustness check on the intersection of solutions scored by every judge, and the bootstrap uncertainty on the alt-test winning rate ω .

Judge-swapped leaderboard. On the cross-judge Benchathon subset (every leaderboard system covered; per-system n in Table 20), each of the five cross-validation judges re-scores the same generations as the primary judge. Per-system n is 15, except Llama-4 ($n=14$), where one generation lacks a parseable score from one judge. Table 19 reports the system-level rank agreement of each judge’s per-system means against the primary judge’s on the identical subset, together with each judge’s mean pass rate at the fixed 50-point threshold, raw and after removing that judge’s measured offset against the blind human pool (Table 21).

Strictness shifts pass rates at the fixed threshold; the offset correction largely restores them, and the ordering itself is stable under every judge.

Per-judge calibration. Table 21 reports the per-judge calibration statistics. It gives two quantities, both averaged over the validation solutions. The direct judge-vs-pool offset $\bar{\Delta}_{\text{dir}} = \text{judge} - \text{full}$ is the mean difference between the judge’s score and the full blind-pool mean; it is the number to quote when prose says “the judge sits X raw points above the pool”. The Calderon pool-substitution shift $\bar{\Delta}_{\text{sub}} = \text{sub} - \text{full}$ compares the pool mean after substituting the judge for one blind reviewer against the full-pool mean; replacing one of $k = 3$ raters scales the direct offset down by $1/3$, so $\bar{\Delta}_{\text{sub}} \approx \bar{\Delta}_{\text{dir}}/3$. Both come with a paired t -test. Shapiro-Wilk does not reject normality in all but two cells; in those two the near-zero offset is unchanged under a Wilcoxon signed-rank cross-check, so the parametric test applies throughout.

Per-pass detail. Table 22 reports the same per-judge calibration statistics as Table 21, but with the three Config-A GPT-5.4-mini intra-judge passes broken out individually alongside the baseline single-pass. This gives a $k = 4$ within-judge stochasticity check on the primary judge. The four per-pass offsets lie close together (stdev 1.5 raw points on human-written solutions, 1.7 on LLM-generated ones). That is small compared to the differences between judges (the lenient judges sit +6 to +18 raw points above the pool on LLM-generated solutions) and to the $\approx \$24$ -raw-point within-solution max-min range across the blind reviewers. The same re-runs give the within-judge stochasticity figure referenced in §5.4: across the $k = 3$ Config-A re-runs on the Benchathon LLM generations, the within-cell standard deviation is 3.45 raw points, while the blind reviewers’ within-solution grades span a 23.9-raw-point max-min range. The two are different dispersion measures, but as expected, a single LLM judge is far more self-consistent across re-runs than the human pool is internally.

Per-annotator alt-test detail. Tables 23 and 24 report the full Calderon blind-pool alt-test by judge for the human-written and LLM-generated solutions respectively. Each block’s header states the winning rate ω that block accumulates. The rows below show the five blind reviewers’ per-annotator n_j , ρ_j^f , ρ_j^h , $\bar{d}_j$, one-sided p -value (Wilcoxon

signed-rank throughout, since every $n_j < 30$), and whether the Benjamini-Yekutieli FDR procedure at $q = 0.05$ across $m = 5$ rejected the null. Only annotators flagged ★ contribute to ω .

Intersection robustness check. Opus-4.7’s Benchathon coverage is $n = 28$ on the human-written subset rather than $n = 30$, because Opus-4.7 did not produce a parseable score on two solutions (the provider’s safety filter or a malformed output). Restricting the alt-test to the intersection of solutions every judge scored ($n = 28$) preserves the two-judge headline: GPT-5.4-mini and Opus-4.7 still clear $\omega \geq 0.5$ ($\omega = 0.80$ and 0.60), and no other judge passes. We report the full- n ω in the main tables and use the intersection as a robustness check only.

Bootstrap uncertainty on ω . At $m = 5$ blind reviewers the alt-test’s winning rate moves in steps of 0.20, so its bootstrap distribution under resampling is wide for every judge. Resampling the validation solutions and re-running the alt-test (500 repetitions) yields a 95% percentile CI on each judge’s ω : no lower bound clears $\omega = 0.5$ for any judge on either subset, and the strongest lower bound is the primary judge’s 0.20 on human-written solutions. This is why the body and [Limitations](#) read the two-judge pass as a promising direction rather than an established result.

O Error-analysis mode counts

Table 25 reports the per-mode instance counts and per-tier flag rates behind the Error analysis section for the two *Falllösung* corpora, which share the five flagged modes. Table 26 covers Doctrinal Principles, which uses the three decision-centric modes of its 4-dimension rubric.

P Reasoning-mode ablation

To isolate the reasoning mechanism on open weights, Qwen3-235B-A22B-*Instruct*-2507, the official non-thinking counterpart of the Thinking-2507 variant in the main evaluation, generated the full Benchathon subset (both generation rounds) and the full Doctrinal Principles corpus through the identical pipeline: same provider, prompts, decoding settings, and primary judge. Table 27 pairs the two variants. Both Benchathon rounds enter the paired comparison, so the Thinking column aggregates more generations than the final-run-only

Table 1 row ([§Experimental setup](#)). The Instruct variant is not part of the 12-system leaderboard.

The toggle’s effect is modest relative to tier gaps. Reasoning-on scores +3.5 raw points on Benchathon and +4.1 on Doctrinal Principles, against a gap of roughly 28 raw points from this system to the Benchathon flagship cluster. On the short-form corpus the gain comes at roughly 43× the per-generation cost; on Benchathon the non-thinking variant spends its budget on roughly twice the visible output instead. An explicit reasoning mechanism alone therefore does not close the tier gap on subsumption-based legal reasoning, consistent with the RQ1 observation that the reasoning-enabled Qwen3-235B underperforms its size.

Q Glossary of German terms

German terms of art are kept alongside English glosses because they lack exact English equivalents and appear verbatim in the released German-language prompts and rubric. Table 28 collects every German term used in the paper.

Table 17: Agreement statistics on the Benchathon human-grading validation subset (design: three blind reviewers per solution with one un-blind creator review reported separately). Judge-vs-human rows report Pearson r , Spearman ρ , MAE, and pass/fail Cohen’s κ , split by solution provenance (human-written vs. LLM-generated). Human IRR rows report ICC(2,1) raw, ICC(2, k) raw, and ICC(2, k) grade-points. The pooled row covers all 45 validation solutions, the indented rows are per-solution-type sub-pools (15 each). Calderon alt-test rows report the per-annotator winning rate ω and average advantage probability ρ at the stated ε , with Benjamini-Yekutieli FDR at $q = 0.05$ across the m blind reviewers. $\omega \geq 0.5$ is the alt-test pass criterion. Blind-pool rows use the three blind reviewers as reference. Creator-reference rows use the un-blind creator grade as a single-expert reference.

Comparison	Sample	Stat 1	Stat 2	Stat 3
Judge vs. mean-human (human sol.)	$n = 30$	$r = 0.76, \rho = 0.76$	MAE raw 10.8, MAE gp 2.69	$\kappa = 0.60$
Judge vs. mean-human (LLM sol.)	$n = 15$	$r = 0.64, \rho = 0.49$	MAE raw 13.1, MAE gp 2.87	$\kappa = 0.40$
Human IRR pooled ($k=3$ blind, all 45)	$n = 45 \times 3$	ICC(2,1) raw 0.63	ICC(2, k) raw 0.84	ICC(2, k) gp 0.79
traditional human	$n = 15 \times 3$	ICC(2,1) raw 0.71	ICC(2, k) raw 0.88	ICC(2, k) gp 0.83
co-creation human	$n = 15 \times 3$	ICC(2,1) raw 0.46	ICC(2, k) raw 0.71	ICC(2, k) gp 0.65
LLM-generated	$n = 15 \times 3$	ICC(2,1) raw 0.69	ICC(2, k) raw 0.87	ICC(2, k) gp 0.83
Calderon blind-pool alt-test (human sol.)	$m = 5 \times n = 30$	$\omega = 0.60 (\varepsilon = 0.15)$	$\rho = 0.62$	passes
Calderon blind-pool alt-test (LLM sol.)	$m = 5 \times n = 15$	$\omega = 0.00 (\varepsilon = 0.15)$	$\rho = 0.41$	fails
Calderon creator-reference alt-test (human sol., $\varepsilon = 0.20$)	$m = 5 \times n = 30$	$\omega = 0.60 (\varepsilon = 0.20)$	$\rho = 0.55$	passes
Calderon creator-reference alt-test (LLM sol., $\varepsilon = 0.20$)	$m = 5 \times n = - - -$	$\omega = 0.00 (\varepsilon = 0.20)$	$\rho = 0.47$	fails

Table 18: LLM-judge configuration. Temperature and max output tokens are the values captured in the per-call metadata of the Benchathon evaluation runs (primary 2026-05-31, cross-validation 2026-05-21 and -31); the retry budget is the worker-side default of 3, with the observed retry count across each run reported in parentheses. No provider-side snapshot id is returned for these aliases; the model id plus run date fixes the configuration for reproducibility.

Role	Model	Run date	T	Max tok.	Retry budget
Primary judge	gpt-5.4-mini	2026-05-31	1.0	8000	3 (obs. 0)
Cross-validation (Anthropic)	claude-opus-4-7	2026-05-21	1.0	8000	3 (obs. 0)
Cross-validation (Anthropic)	claude-sonnet-4-6	2026-05-31	1.0	8000	3 (obs. 0)
Cross-validation (Google)	gemini-3.1-pro-preview	2026-05-21	1.0	8000	3 (obs. 0)
Cross-validation (DeepSeek)	deepseek-ai/DeepSeek-V4-Pro	2026-05-31	1.0	8000	3 (obs. 0)
Cross-validation (Alibaba)	Qwen/Qwen3.5-397B-A17B	2026-05-31	1.0	8000	3 (obs. 0)

Table 19: System-level rank agreement of each cross-validation judge against the primary judge on the shared Benchathon subset (Spearman ρ / Kendall τ over 12 per-system means), and each judge’s mean pass rate at the fixed $\geq 50/100$ threshold, raw and after removing the judge’s measured offset vs. the blind human pool.

Judge	ρ	τ	Pass (raw)	Pass (offset-corr.)
GPT-5.4-mini (primary)	—	—	62%	62%
Opus-4.7	0.97	0.85	58%	55%
Sonnet-4.6	0.95	0.85	71%	58%
Gemini-3.1-Pro	0.92	0.79	74%	46%
DeepSeek-V4-Pro	0.94	0.84	59%	53%
Qwen3.5-397B	0.92	0.79	88%	58%

Table 20: Judge-swapped Benchathon leaderboard: per-system mean raw score under each judge on the identical generation subset, systems ordered by the primary judge ($\dagger$). Strictness moves the level; the ordering is preserved (rank agreement in Table 19).

System	n	GPT-5.4-mini $\dagger$	Opus-4.7	Sonnet-4.6	Gemini-3.1-Pro	DeepSeek-V4-Pro	Qwen3.5-397B
Opus-4.7	15	69.3	71.7	74.7	82.8	76.5	81.5
Gemini-3.1-Pro	15	68.3	68.4	69.6	81.7	74.6	76.3
GPT-5.4	15	66.6	71.5	72.1	83.7	73.3	82.2
Gemini-3.1-Flash-Lite	15	63.5	63.1	67.1	82.0	68.1	77.3
Sonnet-4.6	15	58.9	66.0	70.3	75.9	74.6	77.4
GPT-5.4-mini	15	58.6	51.1	57.3	60.3	50.0	65.5
DeepSeek-V4-Pro	15	56.1	56.9	62.5	70.1	60.5	71.4
Qwen3.5-122B	15	49.4	42.8	48.5	54.6	49.6	59.6
DeepSeek-V4-Flash	15	49.3	47.8	55.3	62.2	52.9	65.1
Qwen3.6-35B	15	42.4	31.5	44.2	40.2	40.7	54.2
Qwen3-235B	15	41.9	34.4	43.9	46.1	39.6	54.1
Llama-4	14	37.7	29.0	37.0	37.8	31.6	51.8

Table 21: Per-judge calibration on the two blind-graded subsets of the Benchathon validation set, ordered by increasing leniency on human-written solutions. Direct offset $\bar{\Delta}_{\text{dir}}$ is the per-solution paired difference between the judge and the mean of three blind reviewers. $\bar{\Delta}_{\text{sub}}$ is the same quantity after the Calderon §1 pool substitution and is attenuated by $1/3$ at $k = 3$ raters. SW p tests normality of the paired sub-full differences. Six judges: GPT-5.4-mini is the primary judge that anchors RQ1-RQ2; Opus-4.7, Sonnet-4.6, DeepSeek-V4-Pro, Gemini-3.1-Pro and Qwen3.5-397B-A17B are Benchathon-only cross-validation judges.

Configuration	n	$\bar{\Delta}_{\text{dir}}$ (sd)	t (p)	$\bar{\Delta}_{\text{sub}}$ (sd)	SW p	t_{sub} (p)
<i>Human-written</i> : Baseline GPT-5.4-mini (primary)	$n = 30$	-0.01 (sd 13.43)	-0.00 (0.9964)	0.20 (sd 4.99)	0.02	0.22 (0.8279)
Opus-4.7	$n = 28$	-2.06 (sd 13.93)	-0.78 (0.4409)	-0.52 (sd 4.71)	0.78	-0.59 (0.5611)
Sonnet-4.6	$n = 30$	0.62 (sd 12.96)	0.26 (0.7945)	0.41 (sd 4.46)	0.94	0.50 (0.6178)
DeepSeek-V4-Pro	$n = 30$	2.11 (sd 14.89)	0.77 (0.4448)	0.91 (sd 4.69)	0.99	1.06 (0.2987)
Gemini-3.1-Pro	$n = 30$	6.82 (sd 15.25)	2.45 (0.0206)	2.48 (sd 4.79)	0.96	2.83 (0.0084)
Qwen3.5-397B-A17B	$n = 30$	8.72 (sd 15.00)	3.19 (0.0034)	3.11 (sd 4.91)	0.31	3.47 (0.0017)
<i>LLM-generated</i> : Baseline GPT-5.4-mini (primary)	$n = 15$	-0.47 (sd 16.24)	-0.11 (0.9130)	0.83 (sd 7.65)	0.33	0.42 (0.6795)
Opus-4.7	$n = 15$	1.37 (sd 10.95)	0.48 (0.6362)	1.44 (sd 5.62)	0.52	0.99 (0.3368)
Sonnet-4.6	$n = 15$	6.87 (sd 13.25)	2.01 (0.0645)	3.28 (sd 6.62)	0.85	1.92 (0.0757)
DeepSeek-V4-Pro	$n = 15$	5.83 (sd 16.98)	1.33 (0.2045)	2.93 (sd 6.63)	0.48	1.71 (0.1088)
Gemini-3.1-Pro	$n = 15$	18.17 (sd 8.40)	8.37 (0.0000)	7.04 (sd 5.03)	0.89	5.42 (0.0001)
Qwen3.5-397B-A17B	$n = 15$	16.67 (sd 13.47)	4.79 (0.0003)	6.54 (sd 6.03)	0.32	4.20 (0.0009)

Table 22: Within-judge stochasticity detail for the primary GPT-5.4-mini judge. The Baseline row is the single-pass run that anchors RQ1-RQ2. The Config A block adds three intra-judge re-runs ($3\times$) so the four offsets in each block give a $k = 4$ within-judge stochasticity check. The per-pass offset $\bar{\Delta}_{\text{dir}}$ stdev across the four passes is 1.53 raw points on human-written solutions and 1.65 on LLM-generated ones, with max $|\text{pass} - \text{mean}|$ of 2.23 (human) and 2.11 (LLM). Same columns as Table 21.

Configuration	n	$\bar{\Delta}_{\text{dir}}$ (sd)	t (p)	$\bar{\Delta}_{\text{sub}}$ (sd)	SW p	t_{sub} (p)
<i>Human-written</i> : Baseline GPT-5.4-mini ($1\times$, primary)	$n = 30$	-0.01 (sd 13.43)	-0.00 (0.9964)	0.20 (sd 4.99)	0.02	0.22 (0.8279)
Config A GPT-5.4-mini ($3\times$ mean)	$n = 30$	-1.68 (sd 14.43)	-0.64 (0.5278)	-0.36 (sd 5.41)	0.10	-0.36 (0.7202)
pass 1	$n = 30$	-0.88 (sd 18.22)	-0.26 (0.7937)	-0.09 (sd 6.62)	0.02	-0.07 (0.9419)
pass 2	$n = 30$	-0.68 (sd 13.82)	-0.27 (0.7901)	-0.02 (sd 5.31)	0.10	-0.02 (0.9819)
pass 3	$n = 30$	-3.49 (sd 13.57)	-1.41 (0.1692)	-0.96 (sd 5.00)	0.57	-1.05 (0.3010)
<i>LLM-generated</i> : Baseline GPT-5.4-mini ($1\times$, primary)	$n = 15$	-0.47 (sd 16.24)	-0.11 (0.9130)	0.83 (sd 7.65)	0.33	0.42 (0.6795)
Config A GPT-5.4-mini ($3\times$ mean)	$n = 15$	2.34 (sd 14.64)	0.62 (0.5450)	1.77 (sd 7.05)	0.38	0.97 (0.3475)
pass 1	$n = 15$	2.97 (sd 17.50)	0.66 (0.5222)	1.98 (sd 7.85)	0.45	0.98 (0.3455)
pass 2	$n = 15$	1.13 (sd 14.89)	0.29 (0.7724)	1.37 (sd 6.77)	0.12	0.78 (0.4472)
pass 3	$n = 15$	2.93 (sd 14.04)	0.81 (0.4319)	1.97 (sd 7.15)	0.89	1.07 (0.3046)

Table 23: Human-written solutions: Calderon §3 alt-test detail by judge. Each block lists the five blind reviewers in the roster (graders 01, 02, 04, 05, 07) for one judge, with that block’s headline ω stated in its italic header. ρ_j^f = mean over $i \in I_j$ of $\mathbf{1}\{S(f, x_i, j) \geq S(h_j, x_i, j)\}$ (probability the LLM judge aligns with the remaining $H_i \setminus \{j\}$ at least as closely as h_j); ρ_j^h symmetrically. $\bar{d}_j = \rho_j^h - \rho_j^f$ is the per-annotator advantage probability gap. p -value is one-sided Wilcoxon signed-rank against $\bar{d}_j \geq \varepsilon = 0.15$. ★ marks rejection under Benjamini-Yekutieli FDR at $q = 0.05$ across $m = 5$.

	n_j	ρ^f / ρ^h	$\bar{d}$	p	rej.
<i>Baseline GPT-5.4-mini (primary) ($\omega = 0.60$ passes)</i>					
grader_01	14	0.57 / 0.43	−0.143	0.0247	
grader_02	19	0.68 / 0.32	−0.368	0.0008	★
grader_04	15	0.73 / 0.33	−0.400	0.0034	★
grader_05	21	0.62 / 0.38	−0.238	0.0021	★
grader_07	21	0.48 / 0.52	0.048	0.0444	
<i>Opus-4.7 ($\omega = 0.60$ passes)</i>					
grader_01	13	0.54 / 0.46	−0.077	0.0471	
grader_02	17	0.59 / 0.41	−0.176	0.0101	★
grader_04	15	0.60 / 0.40	−0.200	0.0128	★
grader_05	19	0.58 / 0.42	−0.158	0.0080	★
grader_07	20	0.40 / 0.60	0.200	0.1650	
<i>Sonnet-4.6 ($\omega = 0.40$)</i>					
grader_01	14	0.50 / 0.50	0.000	0.0676	
grader_02	19	0.63 / 0.42	−0.211	0.0070	★
grader_04	15	0.53 / 0.47	−0.067	0.0365	
grader_05	21	0.67 / 0.33	−0.333	0.0007	★
grader_07	21	0.52 / 0.48	−0.048	0.0175	
<i>DeepSeek-V4-Pro ($\omega = 0.20$)</i>					
grader_01	14	0.64 / 0.50	−0.143	0.0453	
grader_02	19	0.68 / 0.32	−0.368	0.0008	★
grader_04	15	0.53 / 0.60	0.067	0.1651	
grader_05	21	0.52 / 0.52	0.000	0.0411	
grader_07	21	0.43 / 0.57	0.143	0.1015	
<i>Gemini-3.1-Pro ($\omega = 0.40$)</i>					
grader_01	14	0.71 / 0.36	−0.357	0.0067	★
grader_02	19	0.74 / 0.26	−0.474	0.0003	★
grader_04	15	0.53 / 0.53	0.000	0.0844	
grader_05	21	0.52 / 0.52	0.000	0.0411	
grader_07	21	0.38 / 0.62	0.238	0.2060	
<i>Qwen3.5-397B-A17B ($\omega = 0.00$)</i>					
grader_01	14	0.50 / 0.57	0.071	0.1479	
grader_02	19	0.47 / 0.53	0.053	0.0567	
grader_04	15	0.40 / 0.60	0.200	0.2106	
grader_05	21	0.33 / 0.67	0.333	0.3667	
grader_07	21	0.33 / 0.67	0.333	0.3667	

Table 24: LLM-generated solutions: Calderon §3 alt-test detail by judge. Each block lists the five blind reviewers in the roster (graders 01, 02, 04, 05, 07) for one judge, with that block’s headline ω stated in its italic header. ρ_j^f = mean over $i \in I_j$ of $\mathbf{1}\{S(f, x_i, j) \geq S(h_j, x_i, j)\}$ (probability the LLM judge aligns with the remaining $H_i \setminus \{j\}$ at least as closely as h_j); ρ_j^h symmetrically. $\bar{d}_j = \rho_j^h - \rho_j^f$ is the per-annotator advantage probability gap. p -value is one-sided Wilcoxon signed-rank against $\bar{d}_j \geq \varepsilon = 0.15$. ★ marks rejection under Benjamini-Yekutieli FDR at $q = 0.05$ across $m = 5$.

	n_j	ρ^f / ρ^h	$\bar{d}$	p	rej.
<i>Baseline GPT-5.4-mini (primary) ($\omega = 0.00$)</i>					
grader_01	7	0.29 / 0.71	0.429	0.5938	
grader_02	11	0.09 / 0.91	0.818	0.9790	
grader_04	9	0.67 / 0.33	−0.333	0.0273	
grader_05	9	0.56 / 0.44	−0.111	0.0820	
grader_07	9	0.44 / 0.56	0.111	0.2129	
<i>Opus-4.7 ($\omega = 0.20$)</i>					
grader_01	7	0.29 / 0.71	0.429	0.5938	
grader_02	11	0.55 / 0.55	0.000	0.1392	
grader_04	9	0.89 / 0.11	−0.778	0.0039	★
grader_05	9	0.78 / 0.22	−0.556	0.0098	
grader_07	9	0.67 / 0.44	−0.222	0.0645	
<i>Sonnet-4.6 ($\omega = 0.00$)</i>					
grader_01	7	0.00 / 1.00	1.000	1.0000	
grader_02	11	0.36 / 0.73	0.364	0.5845	
grader_04	9	0.67 / 0.33	−0.333	0.0273	
grader_05	9	0.67 / 0.33	−0.333	0.0273	
grader_07	9	0.56 / 0.44	−0.111	0.0820	
<i>DeepSeek-V4-Pro ($\omega = 0.00$)</i>					
grader_01	7	0.00 / 1.00	1.000	1.0000	
grader_02	11	0.36 / 0.64	0.273	0.3501	
grader_04	9	0.78 / 0.22	−0.556	0.0098	
grader_05	9	0.78 / 0.44	−0.333	0.0371	
grader_07	9	0.33 / 0.67	0.333	0.4551	
<i>Gemini-3.1-Pro ($\omega = 0.00$)</i>					
grader_01	7	0.14 / 0.86	0.714	0.8906	
grader_02	11	0.00 / 1.00	1.000	1.0000	
grader_04	9	0.56 / 0.44	−0.111	0.0820	
grader_05	9	0.22 / 0.78	0.556	0.7520	
grader_07	9	0.44 / 0.56	0.111	0.2129	
<i>Qwen3.5-397B-A17B ($\omega = 0.00$)</i>					
grader_01	7	0.00 / 1.00	1.000	1.0000	
grader_02	11	0.18 / 0.91	0.727	0.9731	
grader_04	9	0.56 / 0.44	−0.111	0.0820	
grader_05	9	0.33 / 0.67	0.333	0.4551	
grader_07	9	0.33 / 0.67	0.333	0.4551	

Table 25: Error-analysis mode counts on the two *Falllösung* corpora under the primary judge: flagged instances per mode, corpus n, and per-tier flag rates.

Corpus	Mode	Count	n	Flagship	Efficiency	Open-weight
Benchathon	Long but shallow	10	180	10.0%	0.0%	4.4%
	Short / missing Subsumtion	2	180	1.7%	0.0%	1.1%
	Wrong-norm framing	11	180	0.0%	3.3%	11.1%
	Correct outcome, weak Subsumtion	0	180	0.0%	0.0%	0.0%
	Wrong outcome, sound Subsumtion	0	180	0.0%	0.0%	0.0%
ZJS	Long but shallow	992	6959	20.2%	4.6%	13.5%
	Short / missing Subsumtion	10	6959	0.2%	0.2%	0.1%
	Wrong-norm framing	541	6959	3.3%	7.1%	11.0%
	Correct outcome, weak Subsumtion	18	6959	0.1%	0.4%	0.3%
	Wrong outcome, sound Subsumtion	16	6959	0.5%	0.1%	0.1%

Table 26: Decision-centric error modes on the Doctrinal Principles corpus (4-dimension rubric) under the primary judge.

Mode	Count	n	Flagship	Efficiency	Open-weight
Wrong decision	1541	6365	19.1%	26.9%	26.7%
Correct decision, weak Subsumtion	175	6365	1.0%	2.1%	4.2%
Wrong decision, sound Subsumtion	222	6365	4.4%	3.6%	2.8%

Table 27: Same-base-model reasoning toggle: Qwen3-235B-A22B Thinking-2507 (reasoning on; main evaluation) vs. Instruct-2507 (reasoning off; ablation) under identical prompts, decoding, provider, and primary judge. Pass = judge raw $\geq 50/100$. Output length and per-generation cost from logged per-call metadata.

Corpus	Variant	n	Judge raw	Pass	Trunc.	Words	\$/gen
Benchathon	Thinking-2507	28	41.3	18%	0	678	\$0.0144
	Instruct-2507	30	37.8	10%	0	1356	\$0.0005
Doctr. Princ.	Thinking-2507	531	67.3	69%	0	72	\$0.0030
	Instruct-2507	531	63.3	67%	0	78	\$0.0001

Term	Gloss
<i>Falllösung</i>	Written case solution: the end-to-end work product that is graded.
<i>Gutachtenstil</i>	Appraisal style: the conclusion follows the reasoning; the register graded here.
<i>Urteilsstil</i>	Judgment style: result first, reasoning after; used in judicial decisions.
<i>Obersatz</i>	Step 1 of the scaffold: the norm hypothesis to be examined.
<i>Definition</i>	Step 2: interpretation of the norm's elements.
<i>Subsumtion</i>	Step 3: anchoring the concrete facts under the norm's elements; broadly, the entire norm-application method.
<i>Ergebnis</i>	Step 4: the resulting legal consequence.
<i>Musterlösung</i>	Expert reference solution against which submissions are graded.
<i>Staatsexamen</i>	State law examination; gate to all regulated legal professions in Germany.
<i>ausreichend</i>	"Sufficient": lowest passing grade (4-6 of 18 points).
<i>vollbefriedigend</i>	"Fully satisfactory" (10-12 points); customary distinction tier (<i>Prädikat</i>).
<i>Referendar</i>	Legal trainee in the two-year practical stage between the state examinations.
<i>Volljurist</i>	Fully qualified jurist (passed both state examinations).
<i>Grundprinzipien</i>	Core doctrinal principles; German name of the Doctrinal Principles corpus.
<i>Bereich(e)</i>	Legal area(s): civil, public, criminal law.
<i>Ja/Nein</i>	Yes/no decision label on Doctrinal Principles items.
<i>Ergebnisrichtigkeit</i>	Rubric dimension: result correctness.
<i>Vollständigkeit</i>	Rubric dimension: issue identification / completeness.
<i>Rechtsgrundlagen</i>	Rubric dimension: legal grounding (norm citation).
<i>Rechtskenntnis</i>	Rubric dimension: doctrinal knowledge.
<i>Schwerpunktsetzung</i>	Rubric dimension: problem depth at the case's focal issues.
<i>Methodischer Stil</i>	Rubric dimension: methodological structure (scaffold execution).
<i>Gliederung</i>	Rubric dimension: organization.
<i>Sprache</i>	Rubric dimension: terminology and language.
<i>Formalia</i>	Rubric dimension: formal correctness.

Table 28: German terms used in the paper, with English glosses.