Replicating LEXam on German First State Exam-Style Essay Cases to Probe the "Statistical" Face of Legal Reasoning

Sebastian Nagl

Abstract

Large language models were trained to predict text, not to apply legal method, and still they write fluent German legal opinions. To measure what that fluency is worth, we replicated the Swiss LEXam benchmark, a lean and carefully validated method for grading machine answers to law exams, on 1,127 German exam cases, changing nothing but the jurisdiction. The replication worked. It reproduced LEXam's ordering of the shared models, cost under a hundred euros, and gave machine and human baselines on German first state exam-style cases under an established protocol. The same data also showed where a single correctness score against one reference solution stops matching German correction practice, which protects defensible alternative solutions and grades the quality of each step of the legal subsumption. We report what the replication showed, why we stopped extending it, how the BenGER benchmark that followed grades German case solutions with a ten-dimension rubric validated against blind human graders, and what an instrument must do to grade German legal work thoroughly.

1 The question

Large language models write legal prose that looks like the real thing. Asked to solve a German practice exam, a current model produces an Obersatz, defines the elements of § 823 Abs. 1 BGB, subsumes the facts and closes with a confident Ergebnis. Yet these systems were never taught legal method. They were trained to predict the next word, which is a statistical process rather than a doctrinal one.

This raises a question that reaches beyond computer science. If a statistical machine can pass our exams, what does that say about the exams? Either the machines only imitate the surface of legal argument, and a careful grader should catch them, or substantial parts of what we grade, such as structure, phrasing and the form of the Gutachtenstil, are regular patterns that a statistical learner can absorb. Both answers are informative. The first locates the legal core of legal work; the second tells us which parts of legal assessment are already statistical in practice.

An empirical answer needs two things that rarely coexist: an established way of scoring machine answers to law exams, and an evaluation that fits the legal system being examined. This paper tells how we tried to have both. We started from LEXam, the Swiss benchmark that set the standard for grading machine answers to law exams, and replicated it as faithfully as we could on German exam cases. The replication succeeded. The same data then showed us where a single correctness score against one reference solution stops matching the way German correctors grade, who must accept defensible alternative solutions and who grade the quality of each step of the subsumption rather than the result alone. At that point we stopped extending the replication and built an instrument of our own, the BenGER benchmark, and we are now taking the next step towards the grading sheets that German examiners use. None of this diminishes LEXam. It is the story of a method that travelled well, and of a jurisdiction whose correction practice asks more of the grader than that method was built to deliver.

2 LEXam, and what it gets right

LEXam is a benchmark of 340 law exams from 116 courses at the University of Zurich, with 7,537 questions in English and German: 2,841 long-form open questions and 4,696 multiple-choice questions.[1] Every open question comes with a reference answer written by the professors and with guidance on the reasoning it expects, whether issue spotting, rule recall or rule application. The evaluation is zero-shot: one prompt, no examples, no other help. For the open questions, an ensemble of three language-model judges scores each answer against the reference answer and keeps the lowest of the three scores, and the authors validated this judge against three legal experts with the alternative annotator test, which asks whether the judge can stand in for one of the experts.[2]

Three design choices make LEXam what it is. It is lean, because one prompt and one judging protocol cover hundreds of exams. It is reproducible, because prompts, data and code are public under permissive licences. And it is strict, because the judge rewards what matches the reference and nothing else. These choices fit the Swiss exams well. Most LEXam questions are focused essay questions that call for a bounded answer: identify the issue, state the rule, apply it to the facts. For such questions, correctness against the reference answer is close to what the corrector measures, and a model that scores well has done the legal work the question asks for. LEXam has become the reference point of the field, and rightly so. When we write "correctness judge" below, we mean LEXam's method.

3 The replication on German exam cases

We ran LEXam's protocol on the corpus of BenGER, our benchmark for subsumption-based legal reasoning in German law, which legal experts curate on an open platform.[3] It holds 1,127 German cases in three tiers, each with an official model solution (Musterlösung). Grundprinzipien are 531 short cases on fundamental legal principles that can be answered with a justified Ja or Nein. Benchathon holds 15 curated exam cases from civil, criminal and public law, together with 220 solutions that human participants wrote at a grading event; it is the one tier where humans and machines answered identical questions under identical conditions. ZJS Fälle are 581 full practice exams (Übungsklausuren) from the Zeitschrift für das Juristische Studium, first state exam-style cases that demand a complete, structured solution, with model solutions that average 32,000 characters.

Every model received LEXam's generation prompt, identical to the original up to the substitution of German for Swiss legal references, with no system prompt and no examples. LEXam's judge prompt, adapted the same way, scored each answer against the model solution; we report the mean from 0 to 100. Fourteen affordable models, among them the three with published Swiss scores, wrote about 23,000 answers, all machine-graded, for 65 to 85 euros in total. Four engineering choices differ from the original and are recorded in a deviations register published with this paper: a single judge model, GPT-5.4-mini, instead of the three-judge ensemble, which reads slightly milder in level; structured JSON output from the judge; an answer cap of 16,000 tokens, which was never binding; and a Ja/Nein track in place of the multiple-choice track, for which our corpus has no data. The judge produced no errors on 23,000 cells.

[1] Yu Fan et al., "LEXam: Benchmarking Legal Reasoning on 340 Law Exams," *Proceedings of the Fourteenth International Conference on Learning Representations*, 2026, https://arxiv.org/abs/2505.12864.

[2] Nitay Calderon et al., "The Alternative Annotator Test for LLM-as-a-Judge: How to Statistically Justify Replacing Human Annotators with LLMs," in *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, ed. Wanxiang Che et al. (Vienna: Association for Computational Linguistics, 2025), 16051–81, https://doi.org/10.18653/v1/2025.acl-long.782.

[3] Sebastian Nagl et al., "BenGER: Benchmarking LLM Systems on Subsumption-Based Legal Reasoning in German Law," in *Findings of the Association for Computational Linguistics: EMNLP 2026*, ed. Hung-Yi Lee et al. (Budapest: Association for Computational Linguistics, 2026), 26759–94, https://doi.org/10.18653/2026.findings-emnlp.1266; Sebastian Nagl and Matthias Grabmair, "BenGER Platform: A Collaborative Web Platform for End-to-End Benchmarking of German Legal Tasks," *Proceedings of the Twenty-First International Conference on Artificial Intelligence and Law* (Singapore), 2026, https://arxiv.org/abs/2604.13583.

Difficulty rises with the length of the task at every price point. The best affordable model per tier reaches 74 of 100 on Grundprinzipien, 57 on Benchathon and 43 on the ZJS Fälle, where the field stretches down to 22. On the Ja/Nein track every model but one beats the 50 per cent chance line, the best at 82. The three models with published Swiss scores allow a direct comparison (Table 1). They score 15 to 22 points lower on the full German Klausuren than on LEXam's Swiss exams, 14 to 16 points lower on Benchathon, and 6 to 11 points higher on the short German principle questions. Same models, same method: the difference lies in the multi-issue Gutachten that German exams demand, not in the language.

Table 1: Scores from 0 to 100 under the LEXam correctness judge. The Swiss values are LEXam's published scores from its stricter three-judge ensemble; our single judge reads milder, so the German Klausur deficit is conservative.

Model	Swiss LEXam	German Klausuren (ZJS Fälle)	German principles (Grundprinzipien)
gpt-4.1-mini	54.6	33.1	60.5
gpt-4o-mini	42.6	27.3	53.6
gemma-3-12b	41.3	21.6	46.9

The human baseline under the identical judge is the result we value most. The 220 Benchathon solutions average 57 of 100, but the pool mixes two working modes. Solutions written with AI assistance (co-creation, n = 156) average 61.6 and beat every model in the affordable field, whose mean is 42.1 and whose best is 57.3. Unaided solutions (n = 63) average 46.7, above the field's mean and below its leaders. The co-creation uplift of 14.9 points under LEXam's judge matches the 15.7 points we later measured under BenGER's rubric, so two independently built graders see the same effect.

The replication also agrees with our own instrument on who is better. Across the eleven models that both pipelines cover on the ZJS Fälle, LEXam's correctness judge and BenGER's ten-dimension rubric produce the same ranking.[4] A single reference-anchored judge, moved to another legal system and to a much longer exam format, sorts models the same way a doctrinal rubric does. For a leaderboard, LEXam's recipe is enough, and a research group can run it on a new jurisdiction within days. That is a credit to how the original was built.

4 Where the correctness score stopped matching German correction

The two pipelines agree on the order and disagree on the level, and that disagreement is where our story turns. The rubric's scores run 1 to 10 points above the correctness judge's, and the gap grows for stronger models. The reason is visible in the rubric's weights: ten of its hundred points reward methodological style and another ten reward structure, terminology and form, credit that a correctness judge excludes by design. For ranking models this is harmless. For telling a German educator what a machine answer is worth, it is a difference in kind, because a German grade is not a degree of agreement with a reference but a decision on the statutory scale from 0 to 18 points, with 4 as the pass mark,[5] reached by grading the craft of the Gutachten alongside its substance.

The second limit lies in the single reference. German examination law gives the candidate an answer latitude (Antwortspielraum): a solution that is defensible (vertretbar) and argued consistently with weighty reasons may not be marked as wrong, even where it departs from the reference solution.[6] An error that

[4] After removing superseded duplicate generations so that every model has one answer per case; with the duplicates left in, one adjacent pair of models swaps places.

[5] § 1 of the Verordnung über eine Noten- und Punkteskala für die erste und zweite juristische Prüfung (JurPrNotSkV) of December 3, 1981, BGBl. I p. 1243.

[6] BVerfG, order of April 17, 1991, 1 BvR 419/81 and 1 BvR 213/83, BVerfGE 84, 34 (55); BVerfG, order of April 17, 1991, 1 BvR 1529/84 and 1 BvR 138/87, BVerfGE 84, 59 (79).

follows from an earlier decision (Folgefehler) counts once and not again at every later step.[7] A judge that scores correctness against one reference answer compares every solution with one path. Where the task admits one path, as most focused Swiss questions do, that is exactly right. Where a full German Klausur forks at a disputed question, the judge cannot tell a wrong turn from a defensible one, and a single score between 0 and 1 gives no way to find out.

Our data show the footprint of this. We scored every answer with standard text-similarity measures, which count words shared with the model solution or compare neural meaning representations,[8] and compared them answer by answer with the rubric grade. On the mid-tier Benchathon cases, word overlap with the model solution tracks the doctrinal grade at a rank correlation of about 0.77; on the full ZJS Fälle it drops to 0.58, and whole-document embedding similarity reaches only 0.23. Split into thirds by similarity and by grade, 36 per cent of the ZJS answers sit in the aligned corners, similar and good or dissimilar and poor, against 22 per cent expected by chance. But 6.5 per cent are dissimilar to the model solution and still well graded, and 5 per cent are highly similar and still fail; one of the latter, written by Claude Opus 4.7, sits in the top similarity third and received 29 of 100 points because it paraphrases the terrain of the solution and omits the decisive doctrinal steps. Similarity is a far cruder instrument than LEXam's judge, but the two share their anchor, and the dissimilar yet well graded answers are the ones a reference-anchored score is most likely to undervalue. BenGER's own judge, which also sees the reference solution, disagreed with human graders most on exactly such cases, results reached by other reasoning paths than the reference solution's.[9]

The third limit is granularity. A German corrector does not grade a case solution as a whole but step by step: whether the issues were spotted, the norms named, their elements defined, the facts subsumed under them, the focal issues treated in depth and the minor ones kept short. The correctness judge pools all of this into one number per answer. Under the rubric, two failure modes dominated, and both are invisible in a single score. Flagship models write long but shallow solutions that run to over 2,000 words and still earn at most half marks on subsumption and problem depth, while smaller open-weight models write well-formed solutions that argue from the wrong norms and lose their points on legal grounding and doctrinal knowledge.[10] Only per-step scores separate the two, and only per-step scores tell an educator what to teach.

None of these limits is a flaw in LEXam. The correctness judge measures substance against a reference, and it measures it well: the model ranking and the human split come out the same under both graders. But it answers a narrower question than a German corrector asks. We considered one further experiment to isolate the judging protocol from the generation prompt and deliberately did not run it, because another number would not have closed the gap we had found. The replication had served its purpose, and we built the German instrument.

[7] VG Braunschweig, judgment of December 20, 2018, 6 A 612/15.

[8] Kishore Papineni et al., "BLEU: A Method for Automatic Evaluation of Machine Translation," in *Proceedings of the 40th Annual Meeting of the Association for Computational Linguistics*, ed. Pierre Isabelle et al. (Philadelphia: Association for Computational Linguistics, 2002), 311–18, https://doi.org/10.3115/1073083.1073135; Chin-Yew Lin, "ROUGE: A Package for Automatic Evaluation of Summaries," in *Text Summarization Branches Out* (Barcelona: Association for Computational Linguistics, 2004), 74–81, https://aclanthology.org/W04-1013/; Satanjeev Banerjee and Alon Lavie, "METEOR: An Automatic Metric for MT Evaluation with Improved Correlation with Human Judgments," in *Proceedings of the ACL Workshop on Intrinsic and Extrinsic Evaluation Measures for Machine Translation and/or Summarization*, ed. Jade Goldstein et al. (Ann Arbor, MI: Association for Computational Linguistics, 2005), 65–72, https://aclanthology.org/W05-0909/; Tianyi Zhang et al., "BERTScore: Evaluating Text Generation with BERT," *Proceedings of the Eighth International Conference on Learning Representations*, 2020, https://openreview.net/forum?id=SkeHuCVFDr; Wei Zhao et al., "MoverScore: Text Generation Evaluating with Contextualized Embeddings and Earth Mover Distance," in *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, ed. Kentaro Inui et al. (Hong Kong: Association for Computational Linguistics, 2019), 563–78, https://doi.org/10.18653/v1/D19-1053.

[9] Nagl et al., "BenGER."

[10] Nagl et al., "BenGER."

5 BenGER: grading the way a German corrector grades

BenGER grades every case solution with a rubric that follows German correction practice.[11] Ten dimensions carry a hundred raw points: result correctness (20), doctrinal knowledge (15), subsumption quality (15), issue identification (10), legal grounding (10), problem depth (10), methodological style (10), organisation (5), terminology (3) and formal correctness (2). Each maps to a step of the Gutachtenstil scaffold, and the judge may award no points for keywords without definition and subsumption. The raw score converts to the 18-point scale by the correctors' usual key, so that 50 raw points mean the pass grade of 4 and 60 to 63 raw points mean 7. The judge, again GPT-5.4-mini, sees the task, the model solution and the answer, and returns a score with a written reason for every dimension.

Because the rubric was written for this benchmark, we validated it against humans. Forty-five Benchathon solutions, human and machine written, were each graded by three blind legal reviewers under the same rubric. The judge tracks the reviewers' mean at a Pearson correlation of 0.76 and agrees with their pass or fail decision at a Cohen's kappa of 0.60, while the reviewers themselves disagree by 23.9 raw points on average between the highest and the lowest grade of a solution, a spread in line with the several grade points that Hufeld measured between German correctors.[12] On the alternative annotator test, two of six judges pass on the human-written solutions.[13] The judge sits inside the cloud of human disagreement, not outside it.

The picture of the German exam under this instrument is this. On the 581 ZJS Fälle the flagship systems pass. Gemini 3.1 Pro averages 61.8 of 100, or 7.3 grade points, with 83 per cent of its exams at "ausreichend" or better, followed by GPT-5.4 at 60.4 (78 per cent) and Claude Opus 4.7 at 58.7 (76 per cent). Below them the field falls away: the weakest open-weight systems average 2.4 to 2.8 grade points and pass 4 to 13 per cent of the exams. The difficulty gradient holds, since Gemini 3.1 Pro falls from 83 on Grundprinzipien to 62 on the full Klausuren. The human baseline splits as before. Unaided participants average 50.0 raw points on Benchathon and co-creation 65.7, which an equivalence test cannot distinguish from the flagship tier's 66.6 to 69.3. Today's best systems write full German Klausuren at the lower end of "befriedigend", and nothing below the flagship tier passes reliably.

Where the judge still diverges from the human graders is instructive: on presentation, and on results reached by reasoning paths other than the reference solution's.[14] The rubric has moved the grader closer to German practice, but general criteria applied against one reference still meet the answer latitude only halfway.

6 From general criteria to the examiner's grading sheet

What a thorough evaluation of German legal work requires can be stated from examination law and from educational measurement, and we have done so in a systematisation of the field.[15] A grade must come with reasons that the candidate can contest. A defensible solution that departs from the reference may not be marked wrong, and a follow-on error may not be counted twice. The score must not move with length, surface style or instructions hidden in the answer. And the examiner, not the system, sets the standard, decides and reconsiders. Measured against these requirements, a correctness score against one reference and a rubric of general criteria are steps on one road rather than its end.

The instrument the requirements point to already exists in German practice: the grading sheet (Bewer-

[11] Nagl et al., "BenGER."

[12] Clemens Hufeld, "Jede Korrektur eine andere Note: Quantitative Untersuchung der Objektivität juristischer Klausurbewertungen," *Zeitschrift für Didaktik der Rechtswissenschaft* 11, no. 1 (2024): 59–83, https://doi.org/10.5771/2196-7261-2024-1-59.

[13] Nagl et al., "BenGER."

[14] Nagl et al., "BenGER."

[15] Sebastian Nagl, "SoK: Grading Open-Text Legal Answers, from Law School Exams to Trustworthy Rubric Judges," arXiv, October 20, 2026, https://doi.org/10.48550/arXiv.2610.19992.

tungsbogen) that some correctors write for one exam, with the expected steps of the solution and the points each earns.[16] Written for a judge, such a sheet becomes an instance-specific rubric. Every step holds a few binary requirements that credit the argument made rather than the bare result. A disputed question that opens or closes later steps gets a path of its own, and the answer is scored on the path the candidate took, never on whichever path scores higher. Each requirement is judged on the correct law or on the candidate's own earlier decision, so that a follow-on error costs once. Cardinal errors and the key from points to grade are fixed before grading and applied by code. The judge quotes the passage that earns each credit, a program checks that the quote exists, and the result reaches the examiner as a proposal with flags that call for a human reading.[17] A companion study, announced in that systematisation, builds this judge for German law exams.

This is where the replication led us. LEXam showed that a lean, reference-anchored judge can be validated and can travel. German correction asked for more: credit for alternatives, one penalty per error, a reason per point, and an examiner who decides.

7 What this means

For legal education, fluent structure is now cheap. Any student and any machine can produce a clean Gutachten scaffold, and grading schemes that award substantial credit for form will overvalue machine-assisted work. The discriminating signal has moved into issue spotting, the weighting of problems and the decisive doctrinal step, and examination design should follow it there.

For the question in our title, part of legal reasoning is demonstrably statistical. The aligned majority in our similarity analysis shows that a substantial share of what German grading rewards is pattern-regular; machines reproduce it and overlap measures can grade it. The mismatch minority localises the rest: the decisive step taken or missed, the weighting of arguments, the recognition of what the case turns on. That remainder is identifiable, teachable method, and the upper grade bands it guards remain human territory for now.

For benchmarking, legal reasoning does not need one benchmark to win. It needs a common methodological foundation and jurisdiction-specific instruments built on it. LEXam supplies the foundation, a validated grading protocol that travels. BenGER supplies a German instrument, and the grading sheet will supply the next one. The level disagreements between the two pipelines show that whoever configures the grader decides how much fluent form counts, and with it what counts as legal quality. If such systems enter legal education, that configuration deserves the transparency and contestability we demand of human examination regimes.

8 Method notes and limitations

Prompts, configurations, the deviations register and the analysis code are published with the paper, the corpus is the published BenGER dataset, and the runs live on the open platform, where the grading pipeline is inspectable end to end.[18] Five limitations remain. The graders are themselves language models; LEXam validated its design against human experts and BenGER against blind human reviewers, but machine judges still inherit biases, plausibly including a mild leniency towards machine-typical prose. We used one judge model rather than LEXam's three-judge ensemble, so the published Swiss scores are comparable in ranking and slightly stricter in level, which makes the Klausur deficit in Table 1 conservative. The

[16] Martin Heidebach, "Prüfen im rechtswissenschaftlichen Studium: Die Korrektur juristischer Hausarbeiten anhand eines verbindlichen Bewertungseinheiten-Systems," *Zeitschrift für Didaktik der Rechtswissenschaft* 2, no. 3 (2015): 205–14, https://doi.org/10.5771/2196-7261-2015-3-205; Michael B. Strecker et al., "KI-Unterstützung und Rohpunkteschemata: Die Zukunft der juristischen Klausurkorrektur?" *Ordnung der Wissenschaft*, no. 1 (2026): 1–34, https://ordnungderwissenschaft.de/2025/ki-unterstuetzung-und-rohpunkteschemata-die-zukunft-der-juristischen-klausurkorrektur/.

[17] Nagl, "SoK."

[18] Nagl et al., "BenGER"; Nagl and Grabmair, "BenGER Platform."

flagship models appear only in the rubric pipeline, so the comparison of humans with flagships under LEXam's judge remains open. Coverage has gaps: one model lacks 17 per cent of its ZJS cells, judge zeros pool genuine failures with a rare unparseable-answer default, and one provider refused a criminal-law fact pattern on content-policy grounds, itself a finding about general-purpose systems on legal material. And fifteen cases with 220 human solutions in two working modes are a small, if clean, sample.

AI usage disclosure

Claude Fable 5 and Claude Fable 5.1 (Anthropic) assisted with the coding work behind this study, including the run configuration tooling and the analysis scripts, and with the wording of parts of this text and of its revision for this volume. The experimental design, the interpretation of the results and the text of this paper are the author's own responsibility. All reported numbers were checked against the underlying data.

References

Banerjee, Satanjeev, and Alon Lavie. "METEOR: An Automatic Metric for MT Evaluation with Improved Correlation with Human Judgments." In *Proceedings of the ACL Workshop on Intrinsic and Extrinsic Evaluation Measures for Machine Translation and/or Summarization*, edited by Jade Goldstein, Alon Lavie, Chin-Yew Lin, and Clare Voss, 65–72. Ann Arbor, MI: Association for Computational Linguistics, 2005. https://aclanthology.org/W05-0909/.

Calderon, Nitay, Roi Reichart, and Rotem Dror. "The Alternative Annotator Test for LLM-as-a-Judge: How to Statistically Justify Replacing Human Annotators with LLMs." In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, edited by Wanxiang Che, Joyce Nabende, Ekaterina Shutova, and Mohammad Taher Pilehvar, 16051–81. Vienna: Association for Computational Linguistics, 2025. https://doi.org/10.18653/v1/2025.acl-long.782.

Fan, Yu, Jingwei Ni, Jakob Merane, et al. "LEXam: Benchmarking Legal Reasoning on 340 Law Exams." *Proceedings of the Fourteenth International Conference on Learning Representations*, 2026. https://arxiv.org/abs/2505.12864.

Heidebach, Martin. "Prüfen im rechtswissenschaftlichen Studium: Die Korrektur juristischer Hausarbeiten anhand eines verbindlichen Bewertungseinheiten-Systems." *Zeitschrift für Didaktik der Rechtswissenschaft* 2, no. 3 (2015): 205–14. https://doi.org/10.5771/2196-7261-2015-3-205.

Hufeld, Clemens. "Jede Korrektur eine andere Note: Quantitative Untersuchung der Objektivität juristischer Klausurbewertungen." *Zeitschrift für Didaktik der Rechtswissenschaft* 11, no. 1 (2024): 59–83. https://doi.org/10.5771/2196-7261-2024-1-59.

Lin, Chin-Yew. "ROUGE: A Package for Automatic Evaluation of Summaries." In *Text Summarization Branches Out*, 74–81. Barcelona: Association for Computational Linguistics, 2004. https://aclanthology.org/W04-1013/.

Nagl, Sebastian. "SoK: Grading Open-Text Legal Answers, from Law School Exams to Trustworthy Rubric Judges." arXiv, October 20, 2026. https://doi.org/10.48550/arXiv.2610.19992.

Nagl, Sebastian, and Matthias Grabmair. "BenGER Platform: A Collaborative Web Platform for End-to-End Benchmarking of German Legal Tasks." *Proceedings of the Twenty-First International Conference on Artificial Intelligence and Law* (Singapore), 2026. https://arxiv.org/abs/2604.13583.

Nagl, Sebastian, Ann-Kristin Mayrhofer, Martin Heidebach, et al. "BenGER: Benchmarking LLM Systems on Subsumption-Based Legal Reasoning in German Law." In *Findings of the Association for Computational Linguistics: EMNLP 2026*, edited by Hung-Yi Lee, Sunipa Dev, Jessy Li, and Desmond Elliott, 26759–94. Budapest: Association for Computational Linguistics, 2026. https://doi.org/10.18653/2026.findings-emnlp.1266.

Papineni, Kishore, Salim Roukos, Todd Ward, and Wei-Jing Zhu. "BLEU: A Method for Automatic Evaluation of Machine Translation." In *Proceedings of the 40th Annual Meeting of the Association for Computational Linguistics*, edited by Pierre Isabelle, Eugene Charniak, and Dekang Lin, 311–18. Philadelphia: Association for Computational Linguistics, 2002. https://doi.org/10.3115/1073083.1073135.

Strecker, Michael B., Susanne Hähnchen, Martin Heidebach, et al. "KI-Unterstützung und Rohpunkteschemata: Die Zukunft der juristischen Klausurkorrektur?" *Ordnung der Wissenschaft*, no. 1 (2026): 1–34. https://ordnungderwissenschaft.de/2025/ki-unterstuetzung-und-rohpunkteschemata-die-zukunft-der-juristischen-klausurkorrektur/.

Zhang, Tianyi, Varsha Kishore, Felix Wu, Kilian Q. Weinberger, and Yoav Artzi. "BERTScore: Evaluating Text Generation with BERT." *Proceedings of the Eighth International Conference on Learning Representations*, 2020. https://openreview.net/forum?id=SkeHuCVFDr.

Zhao, Wei, Maxime Peyrard, Fei Liu, Yang Gao, Christian M. Meyer, and Steffen Eger. "MoverScore: Text Generation Evaluating with Contextualized Embeddings and Earth Mover Distance." In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, edited by Kentaro Inui, Jing Jiang, Vincent Ng, and Xiaojun Wan, 563–78. Hong Kong: Association for Computational Linguistics, 2019. https://doi.org/10.18653/v1/D19-1053.